

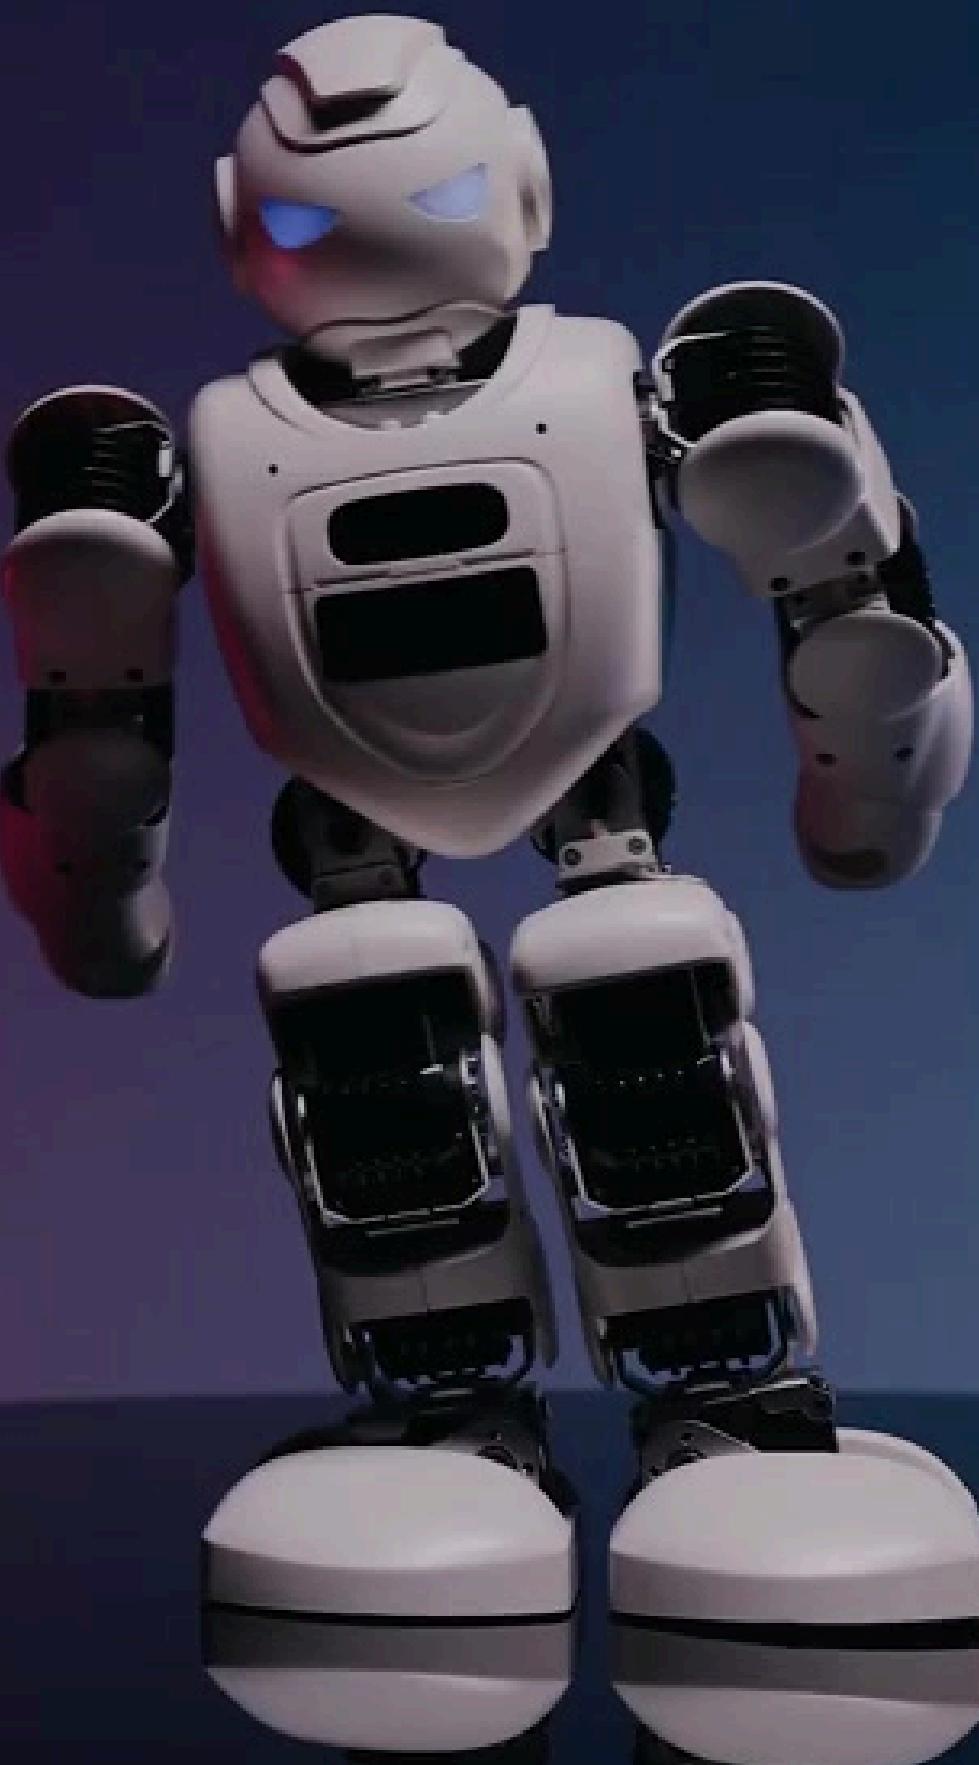


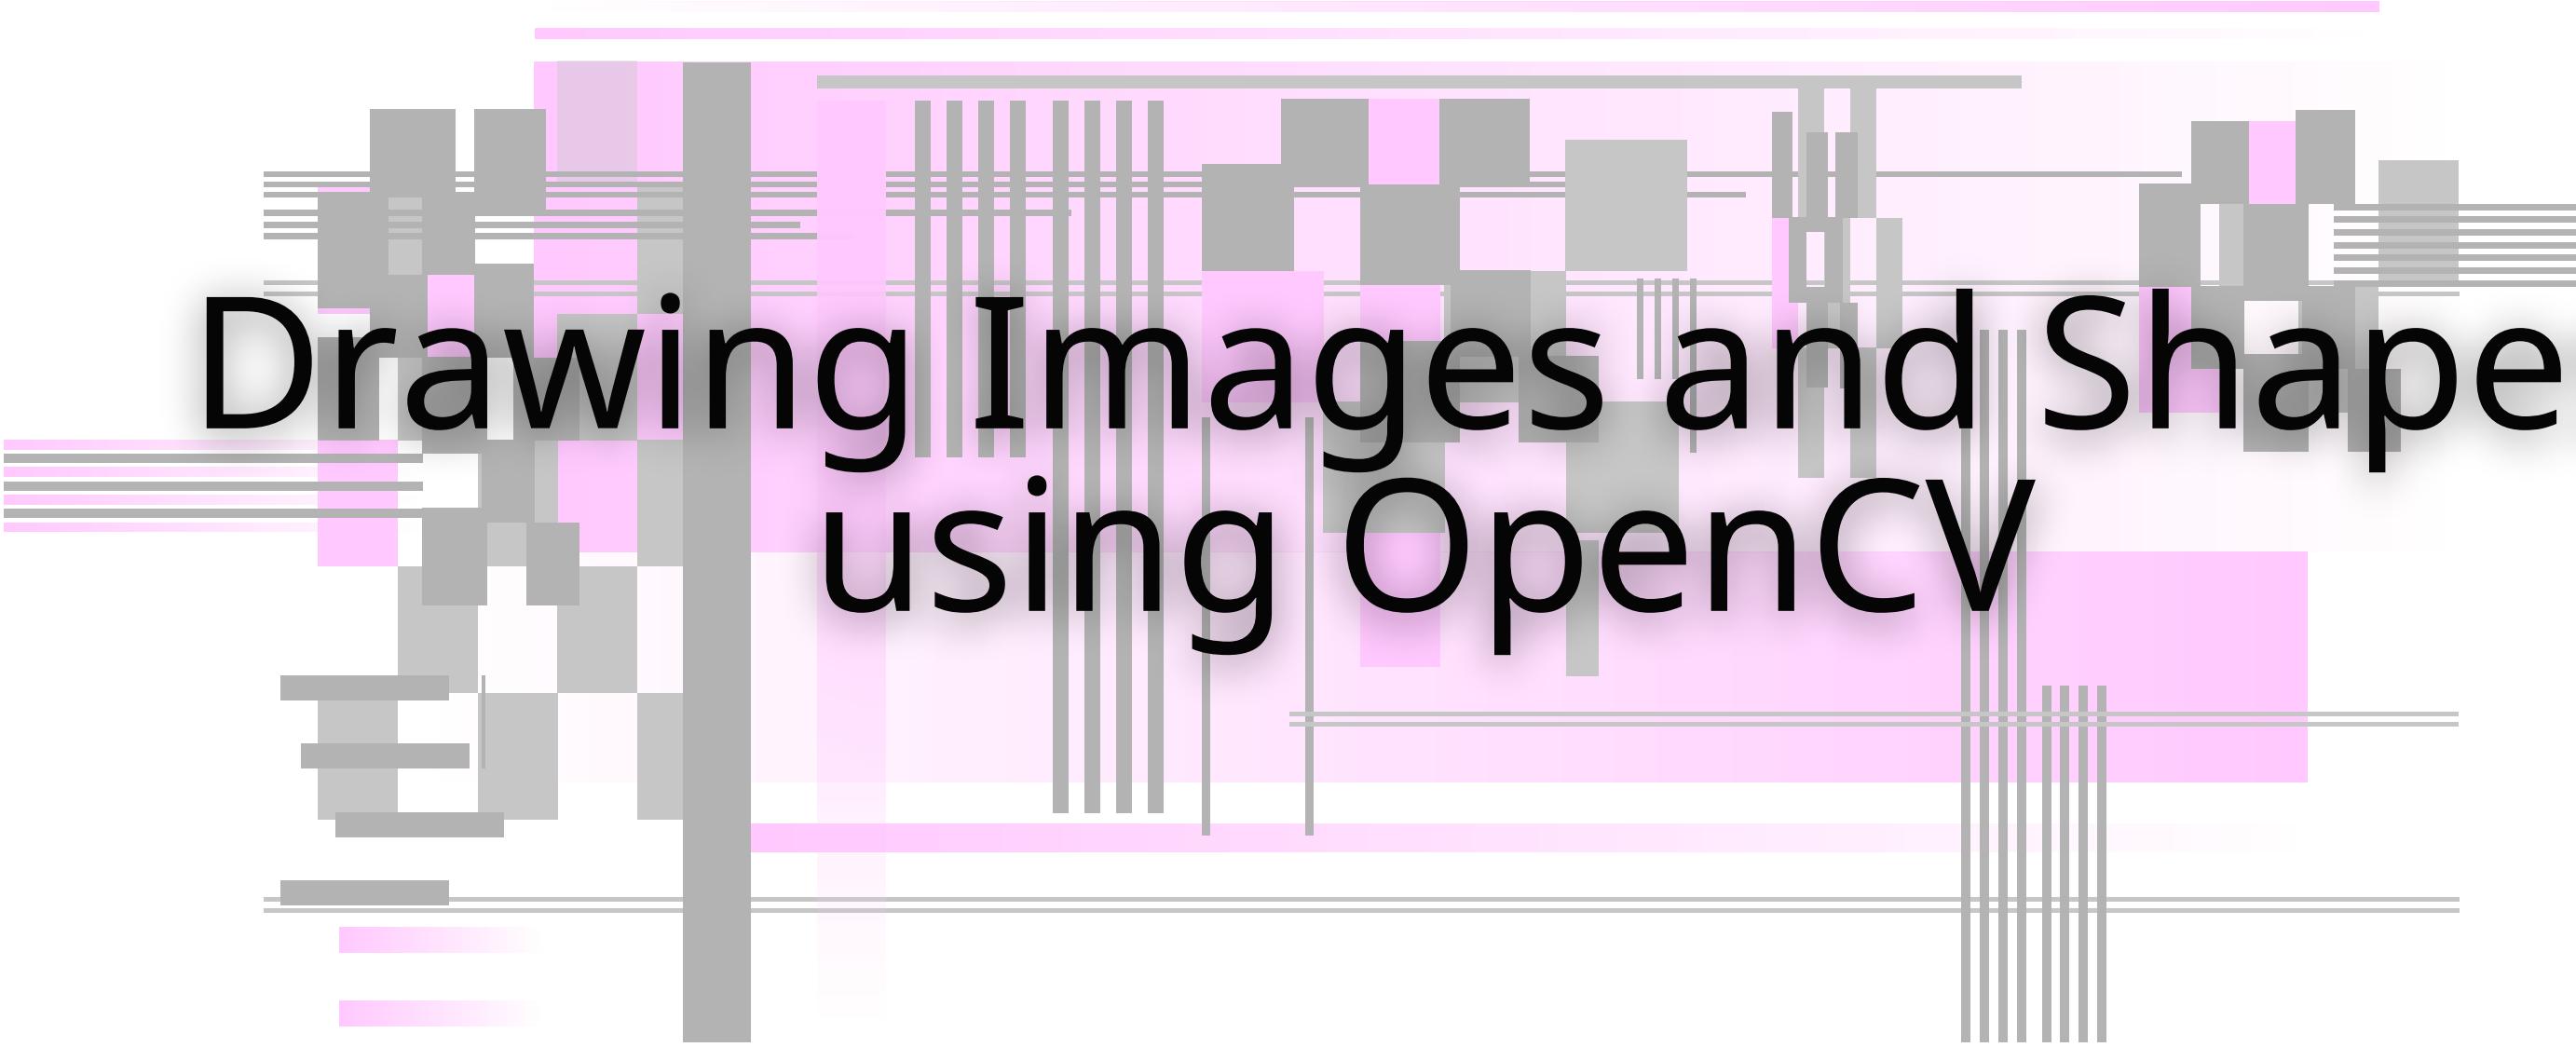
WEEK 3

DEEP LEARNING FOR COMPUTER VISION

UNLIMITED

Presented by **Asst. Prof. Dr. Tuchsanai Ploysuwan**





Drawing Images and Shapes using OpenCV

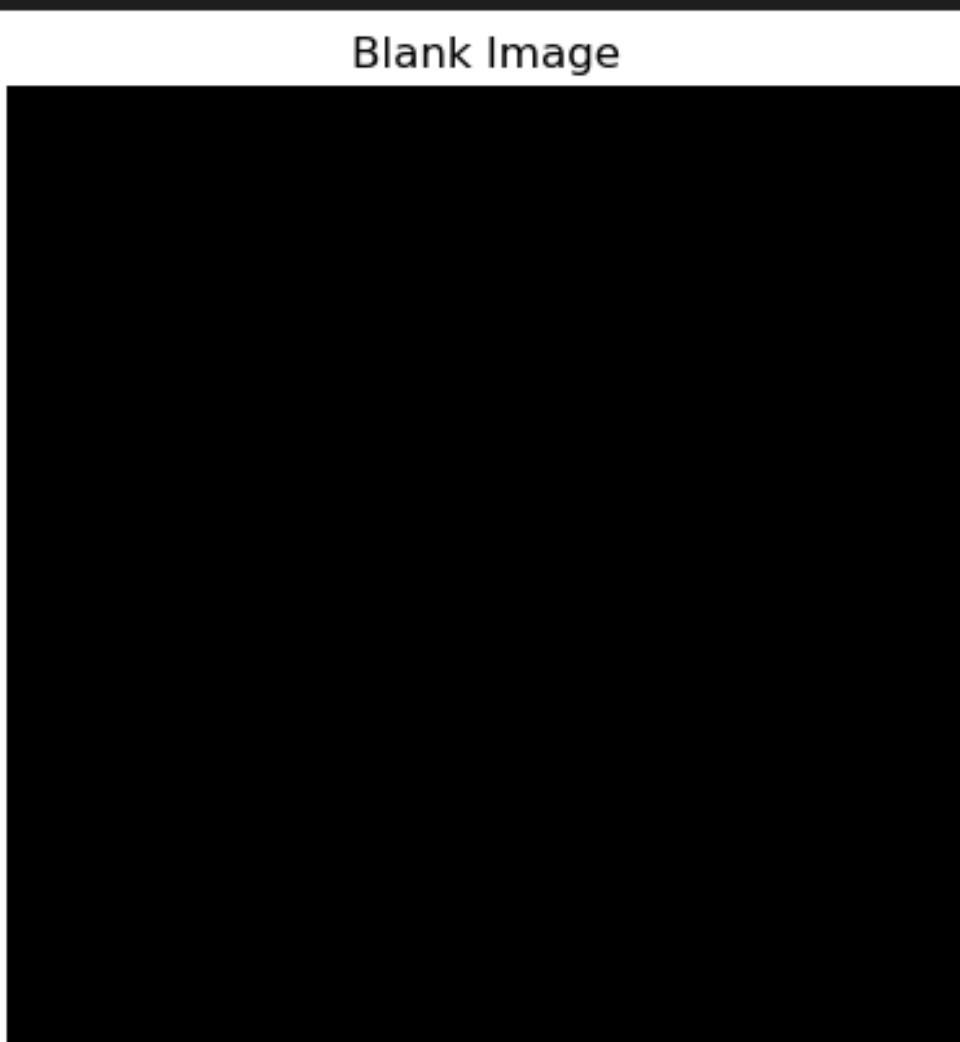
```
import cv2
import numpy as np
import matplotlib.pyplot as plt

def show_image(image, title="Image"):
    # Function to display image using matplotlib
    plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB))
    plt.title(title)
    plt.axis('off')
    plt.show()
```

✓ 0.3s

```
# Create a blank image - a 500x500 image with 3 color channels (BGR) filled with black
blank_image = np.zeros((500, 500, 3), dtype='uint8')
show_image(blank_image, title="Blank Image")
```

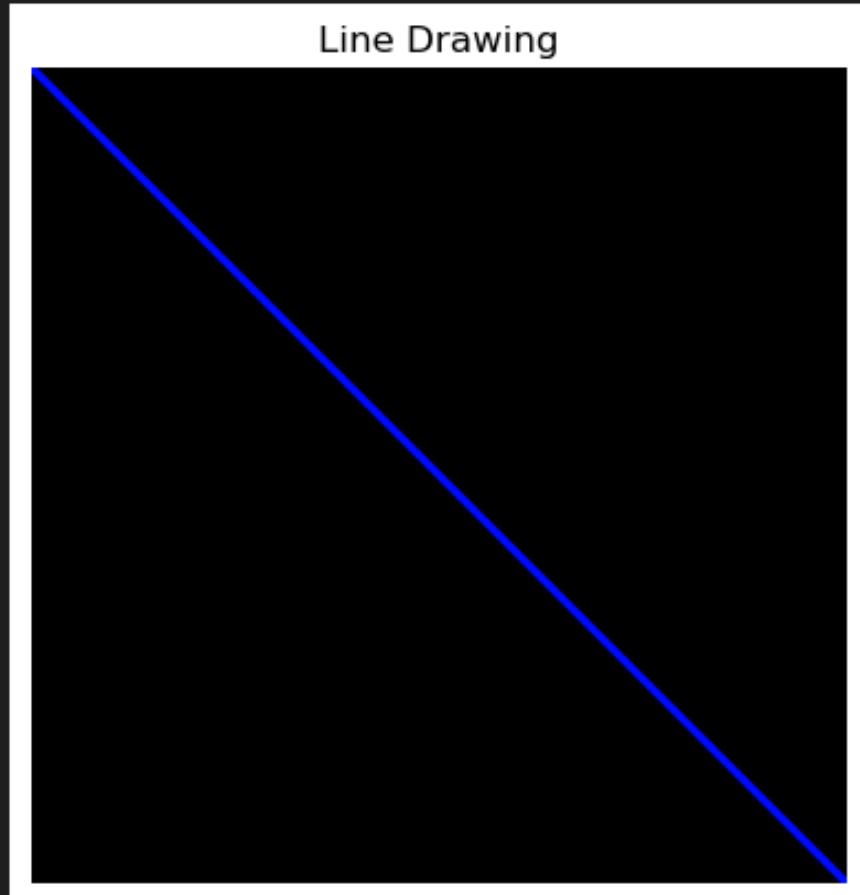
✓ 0.0s



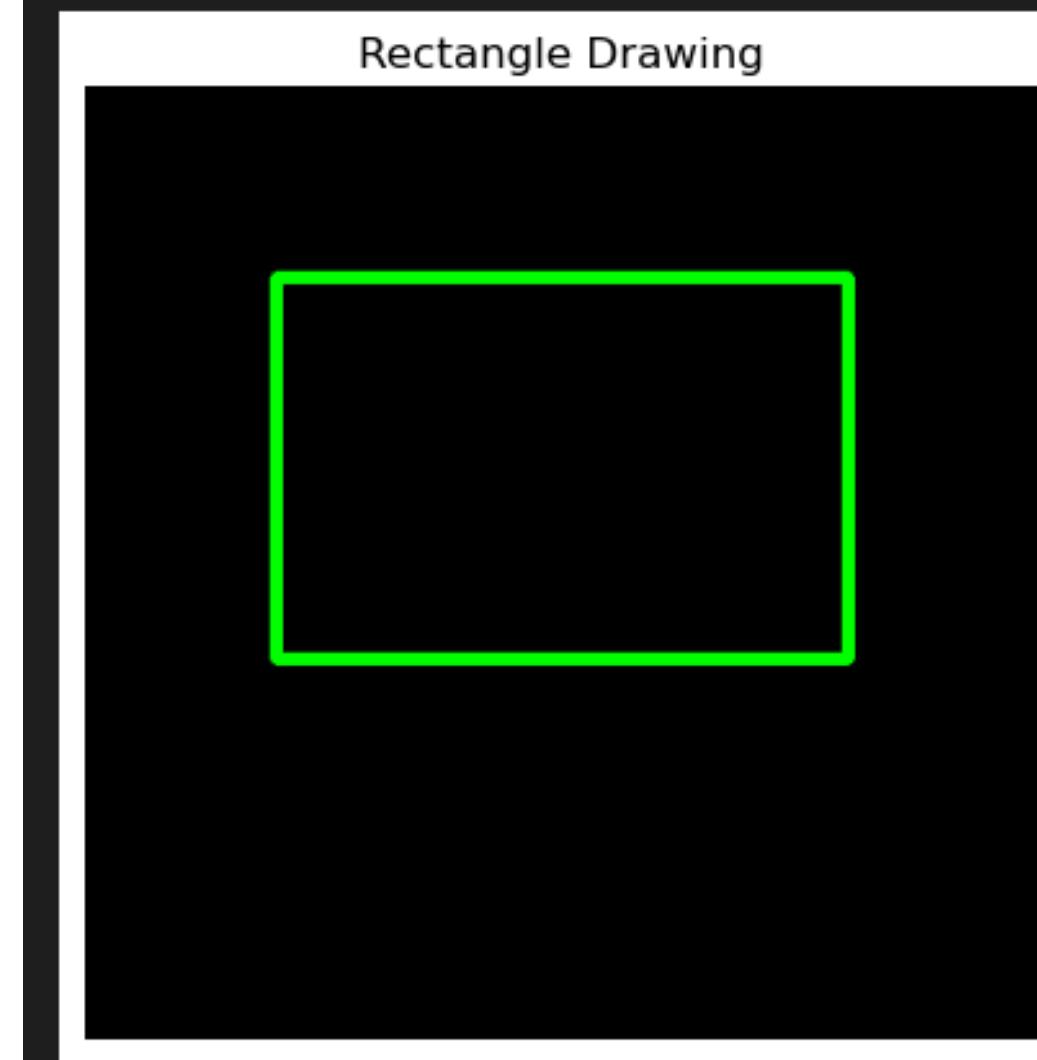
Step 2: Drawing a Line

Let's draw a line from the top-left corner to the bottom-right corner.

```
# Draw a blue line from top-left to bottom-right
line_image = blank_image.copy()
cv2.line(line_image, (0, 0), (500, 500), (255, 0, 0), thickness=3)
show_image(line_image, title="Line Drawing")
```



```
# Draw a green rectangle
rectangle_image = blank_image.copy()
cv2.rectangle(rectangle_image, (100, 100), (400, 300), (0, 255, 0), thickness=5)
show_image(rectangle_image, title="Rectangle Drawing")
```



Explanation of `cv2.line` Function:

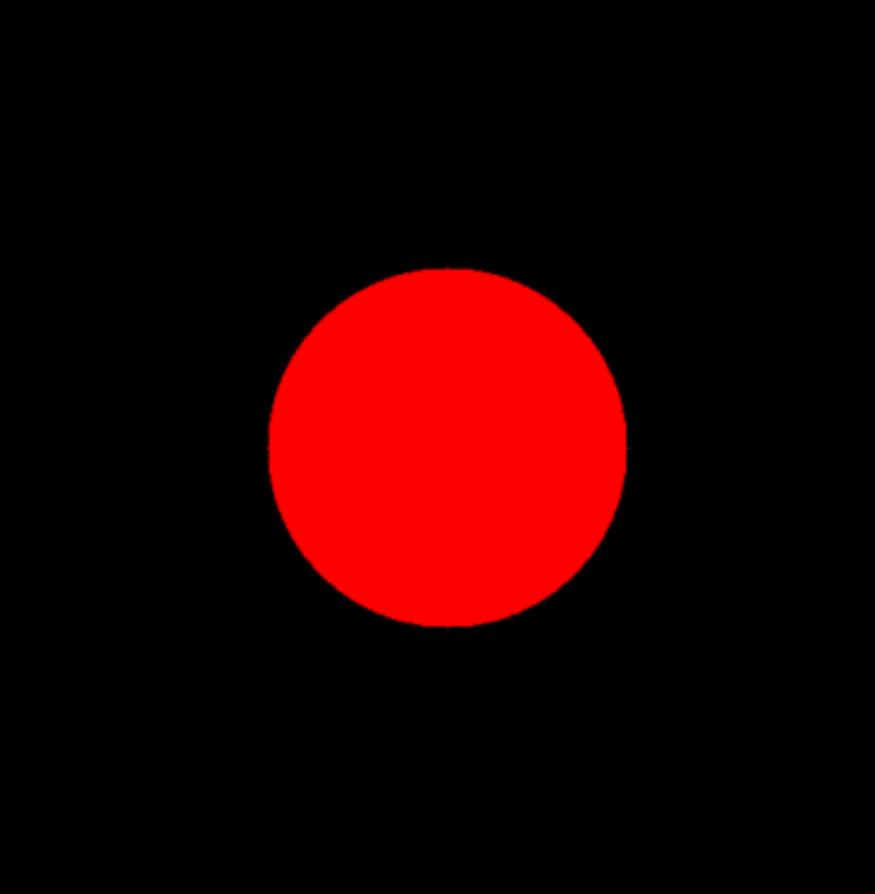
- **Function:** `cv2.line(image, start_point, end_point, color, thickness)`
- **Parameters:**
 - `image`: The image on which to draw.
 - `start_point`: Tuple representing the starting point of the line (x, y).
 - `end_point`: Tuple representing the ending point of the line (x, y).
 - `color`: BGR tuple for the color of the line (e.g., `(255, 0, 0)` for blue).
 - `thickness`: Thickness of the line in pixels.

Explanation of `cv2.rectangle` Function:

- **Function:** `cv2.rectangle(image, top_left, bottom_right, color, thickness)`
- **Parameters:**
 - `image`: The image on which to draw.
 - `top_left`: Tuple representing the top-left corner (x, y).
 - `bottom_right`: Tuple representing the bottom-right corner (x, y).
 - `color`: BGR tuple for the color of the rectangle (e.g., `(0, 255, 0)` for green).
 - `thickness`: Thickness of the rectangle's border. Use `-1` to fill the rectangle.

```
# Draw a red circle
# thickness = -1 fills the circle
circle_image = blank_image.copy()
cv2.circle(circle_image, (250, 250), radius=100, color=(0, 0, 255), thickness=-1)
show_image(circle_image, title="Circle Drawing")
```

Circle Drawing

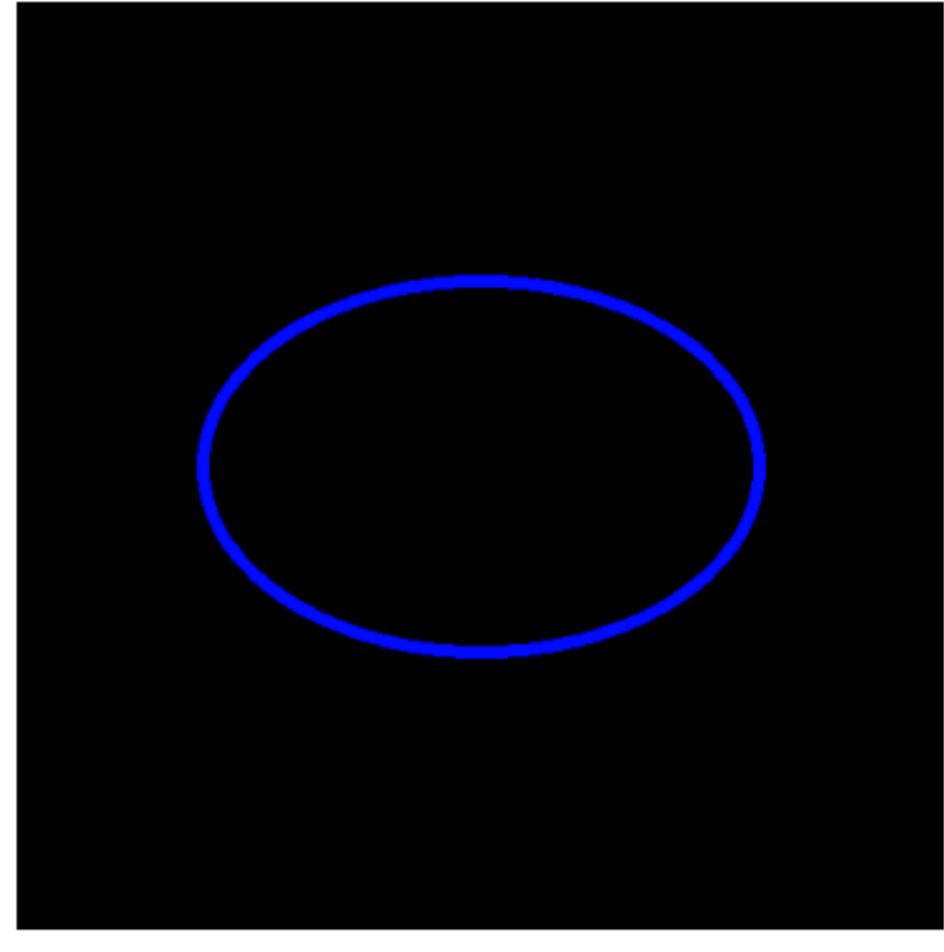


Explanation of `cv2.circle` Function:

- **Function:** `cv2.circle(image, center, radius, color, thickness)`
- **Parameters:**
 - `image`: The image on which to draw.
 - `center`: Tuple representing the center of the circle (x, y).
 - `radius`: Radius of the circle.
 - `color`: BGR tuple for the color of the circle (e.g., `(0, 0, 255)` for red).
 - `thickness`: Thickness of the circle's border. Use `-1` to fill the circle.

```
# Draw an ellipse
ellipse_image = blank_image.copy()
cv2.ellipse(ellipse_image, (250, 250), (150, 100), 0, 0, 360, (255, 0, 0), thickness=5)
show_image(ellipse_image, title="Ellipse Drawing")
```

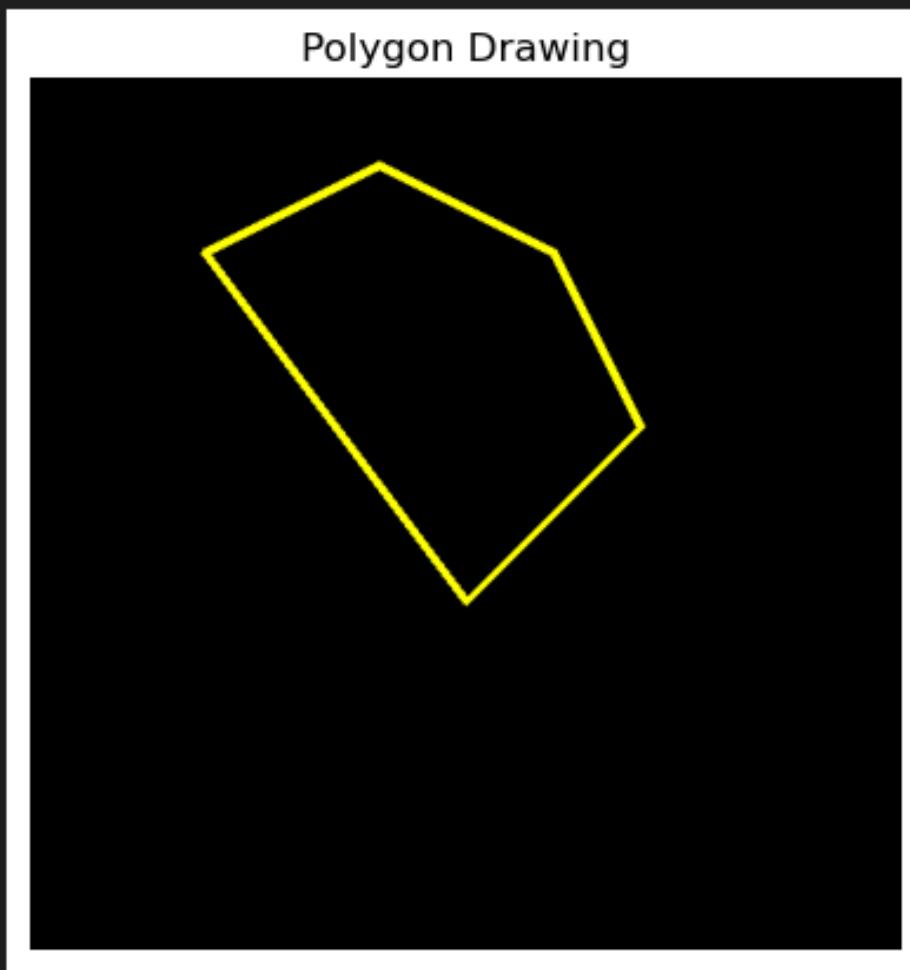
Ellipse Drawing



Explanation of `cv2.ellipse` Function:

- **Function:** `cv2.ellipse(image, center, axes, angle, startAngle, endAngle, color, thickness)`
- **Parameters:**
 - `image`: The image on which to draw.
 - `center`: Tuple representing the center of the ellipse (x, y).
 - `axes`: Tuple representing the lengths of the ellipse axes (major axis, minor axis).
 - `angle`: Angle of rotation of the ellipse in degrees.
 - `startAngle`: Starting angle of the elliptic arc in degrees.
 - `endAngle`: Ending angle of the elliptic arc in degrees.
 - `color`: BGR tuple for the color of the ellipse.
 - `thickness`: Thickness of the ellipse's border.

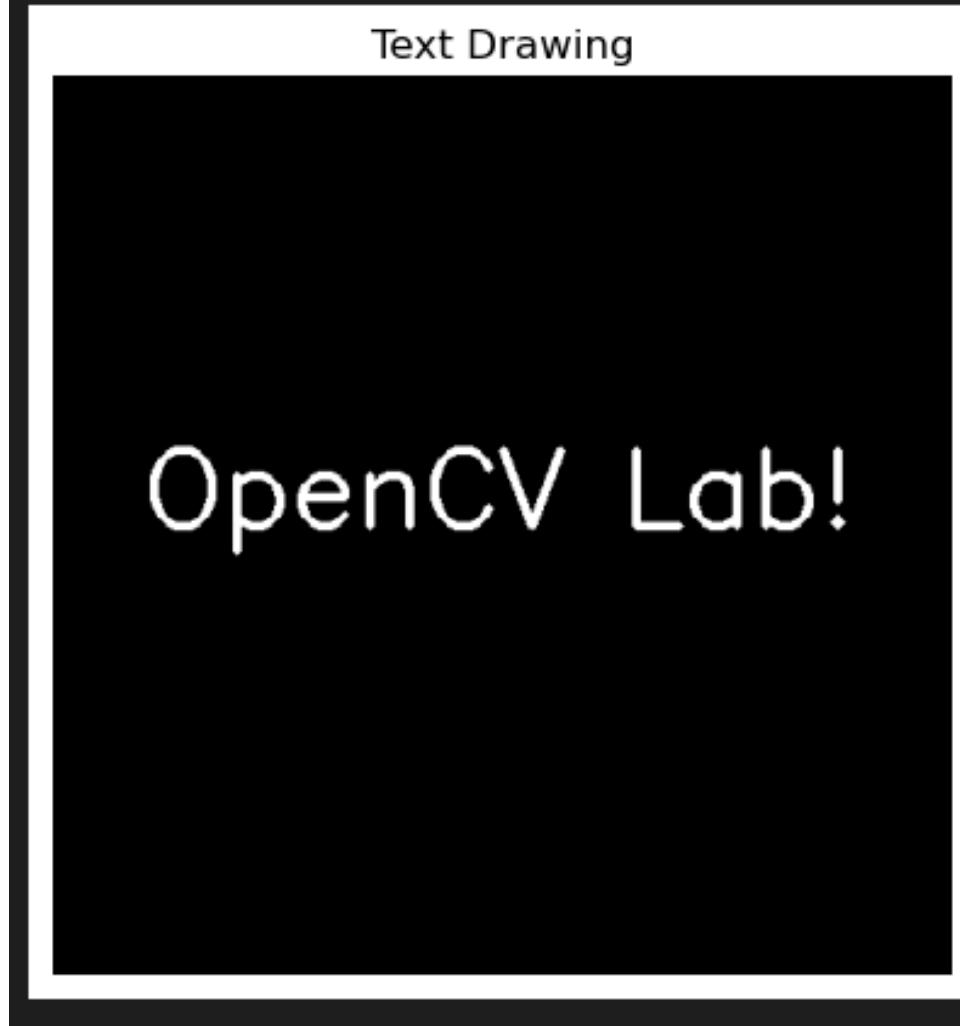
```
# Draw a closed polygon
polygon_image = blank_image.copy()
points = np.array([[100, 100], [200, 50], [300, 100], [350, 200], [250, 300]], np.int32)
points = points.reshape((-1, 1, 2))
cv2.polyline(polygon_image, [points], isClosed=True, color=(0, 255, 255), thickness=3)
show_image(polygon_image, title="Polygon Drawing")
```



Explanation of `cv2.polyline` Function:

- **Function:** `cv2.polyline(image, points, isClosed, color, thickness)`
- **Parameters:**
 - `image`: The image on which to draw.
 - `points`: Array of points representing the vertices of the polygon.
 - `isClosed`: Boolean value specifying whether the polygon should be closed or not.
 - `color`: BGR tuple for the color of the polygon.
 - `thickness`: Thickness of the polygon's border.

```
# Add text to the image
text_image = blank_image.copy()
cv2.putText(text_image, 'OpenCV Lab!', (50, 250), cv2.FONT_HERSHEY_SIMPLEX, 2, (255, 255, 255), 3)
show_image(text_image, title="Text Drawing")
```



House Drawing

```
# Draw a simple house
house_image = blank_image.copy()

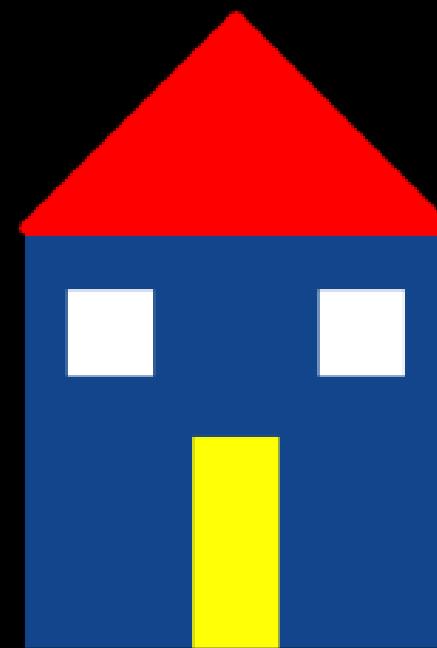
# Draw the main rectangle for the house
cv2.rectangle(house_image, (150, 200), (350, 400), (139, 69, 19), thickness=-1) # Brown filled rectangle

# Draw the roof using a triangle
roof_points = np.array([[150, 200], [350, 200], [250, 100]], np.int32)
roof_points = roof_points.reshape((-1, 1, 2))
cv2.polyline(house_image, [roof_points], isClosed=True, color=(0, 0, 255), thickness=5) # Red outline for the roof
cv2.fillPoly(house_image, [roof_points], color=(0, 0, 255)) # Fill the roof with red color

# Draw the door
cv2.rectangle(house_image, (230, 300), (270, 400), (0, 255, 255), thickness=-1) # Yellow filled rectangle for the door

# Draw windows
cv2.rectangle(house_image, (170, 230), (210, 270), (255, 255, 255), thickness=-1) # White filled rectangle for left window
cv2.rectangle(house_image, (290, 230), (330, 270), (255, 255, 255), thickness=-1) # White filled rectangle for right window

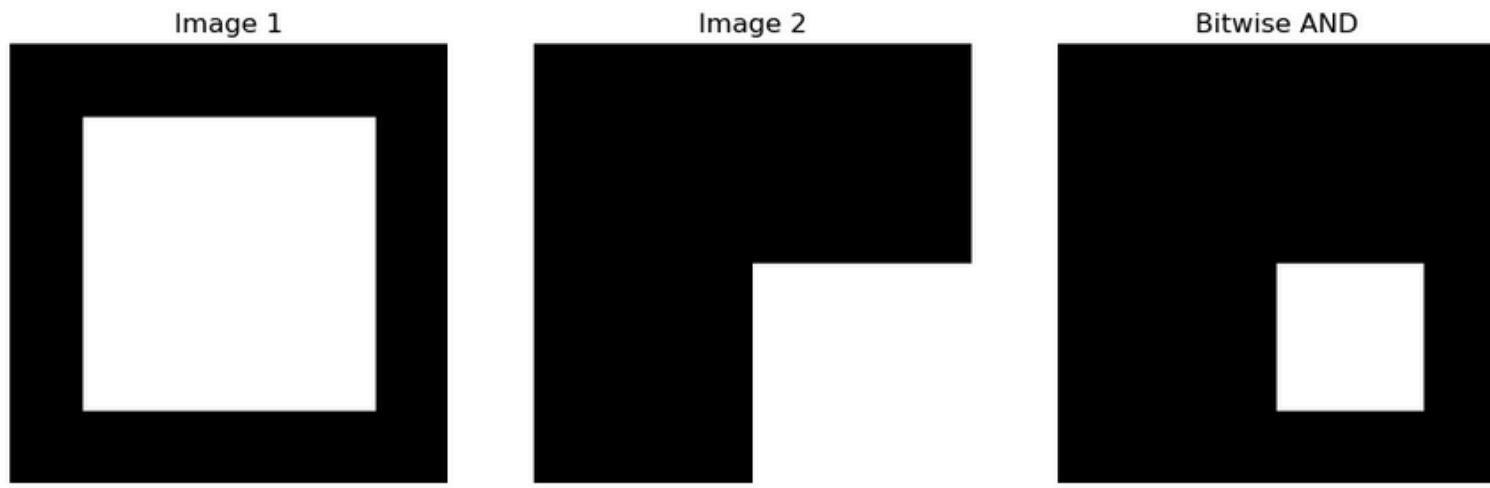
show_image(house_image, title="House Drawing")
```



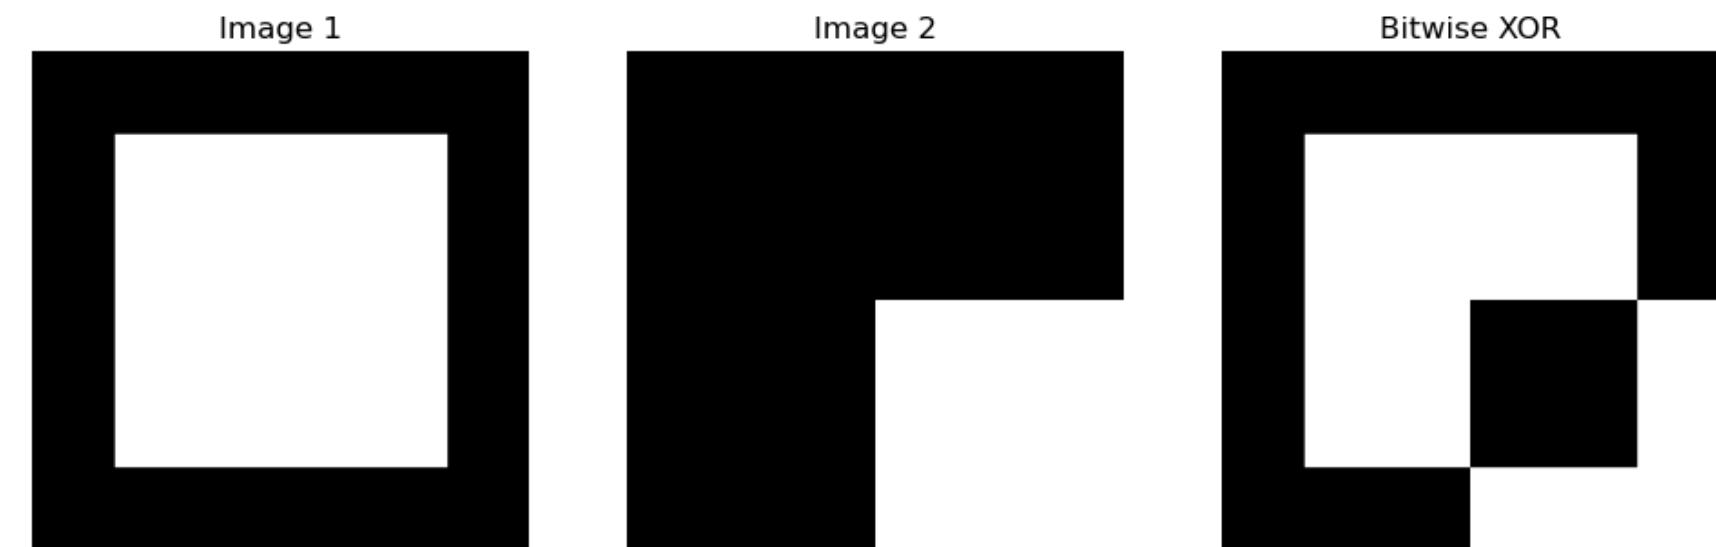
Bitwise Operations



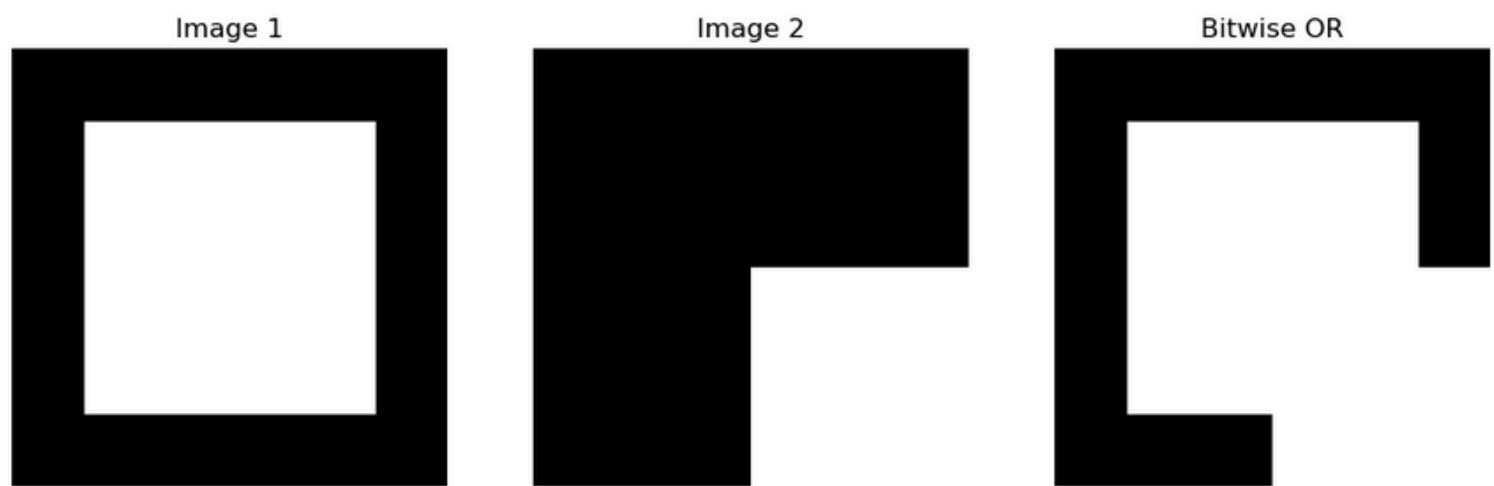
```
# Bitwise AND  
and_result = cv2.bitwise_and(img1, img2)
```



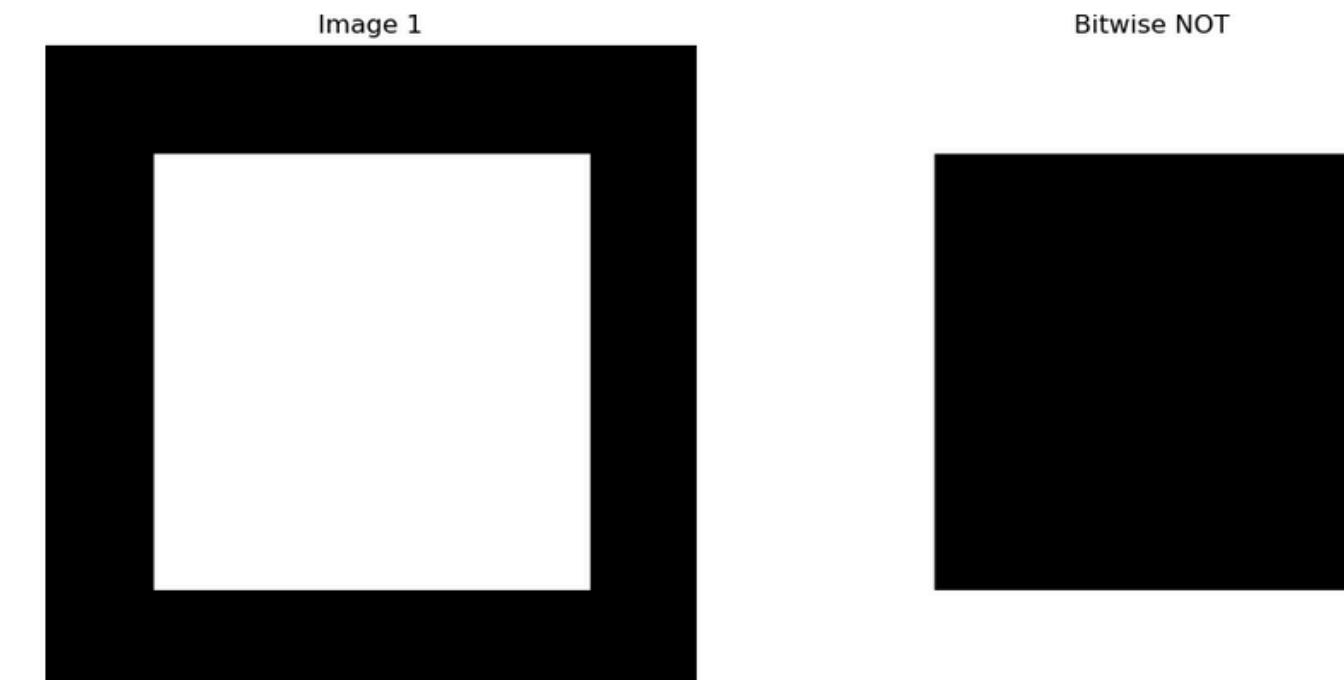
```
# Bitwise XOR  
xor_result = cv2.bitwise_xor(img1, img2)
```



```
# Bitwise OR  
or_result = cv2.bitwise_or(img1, img2)
```



```
# Bitwise NOT  
not_result = cv2.bitwise_not(img1)
```



Practical Example - Masking an Image

```
# Load a sample image
image = cv2.imread(cv2.samples.findFile('./images/Taj_Mahal.jpg'), cv2.IMREAD_GRAYSCALE)

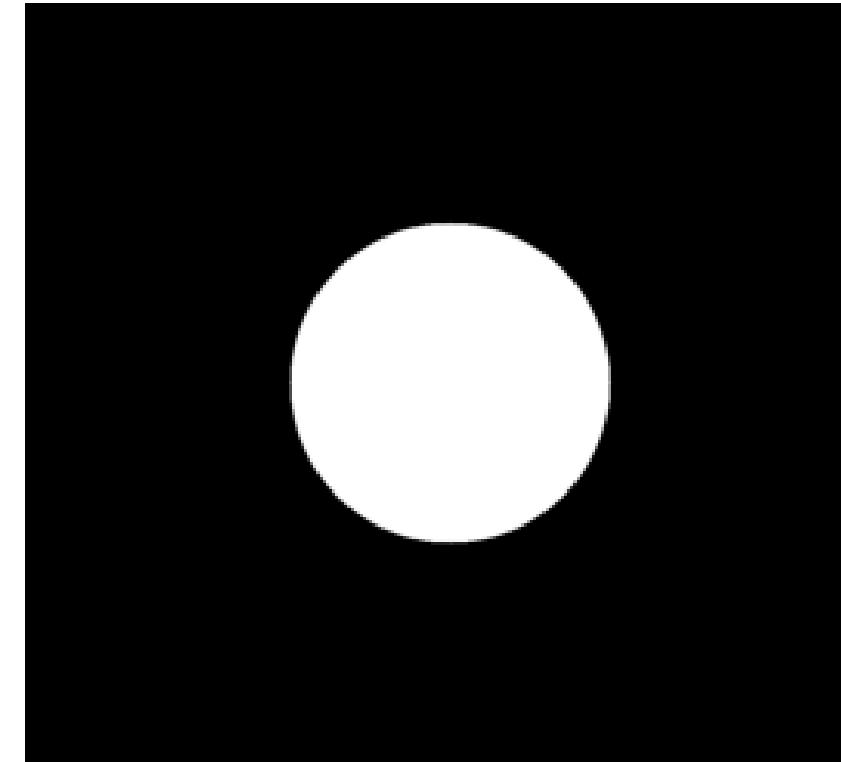
# Create a circular mask
mask = np.zeros_like(image, dtype=np.uint8)
cv2.circle(mask, (mask.shape[1] // 2, mask.shape[0] // 2), 150, 255, -1)

# Apply the mask using Bitwise AND
masked_image = cv2.bitwise_and(image, image, mask=mask)
```

Original Image



Mask



Masked Image



Morphological Operations

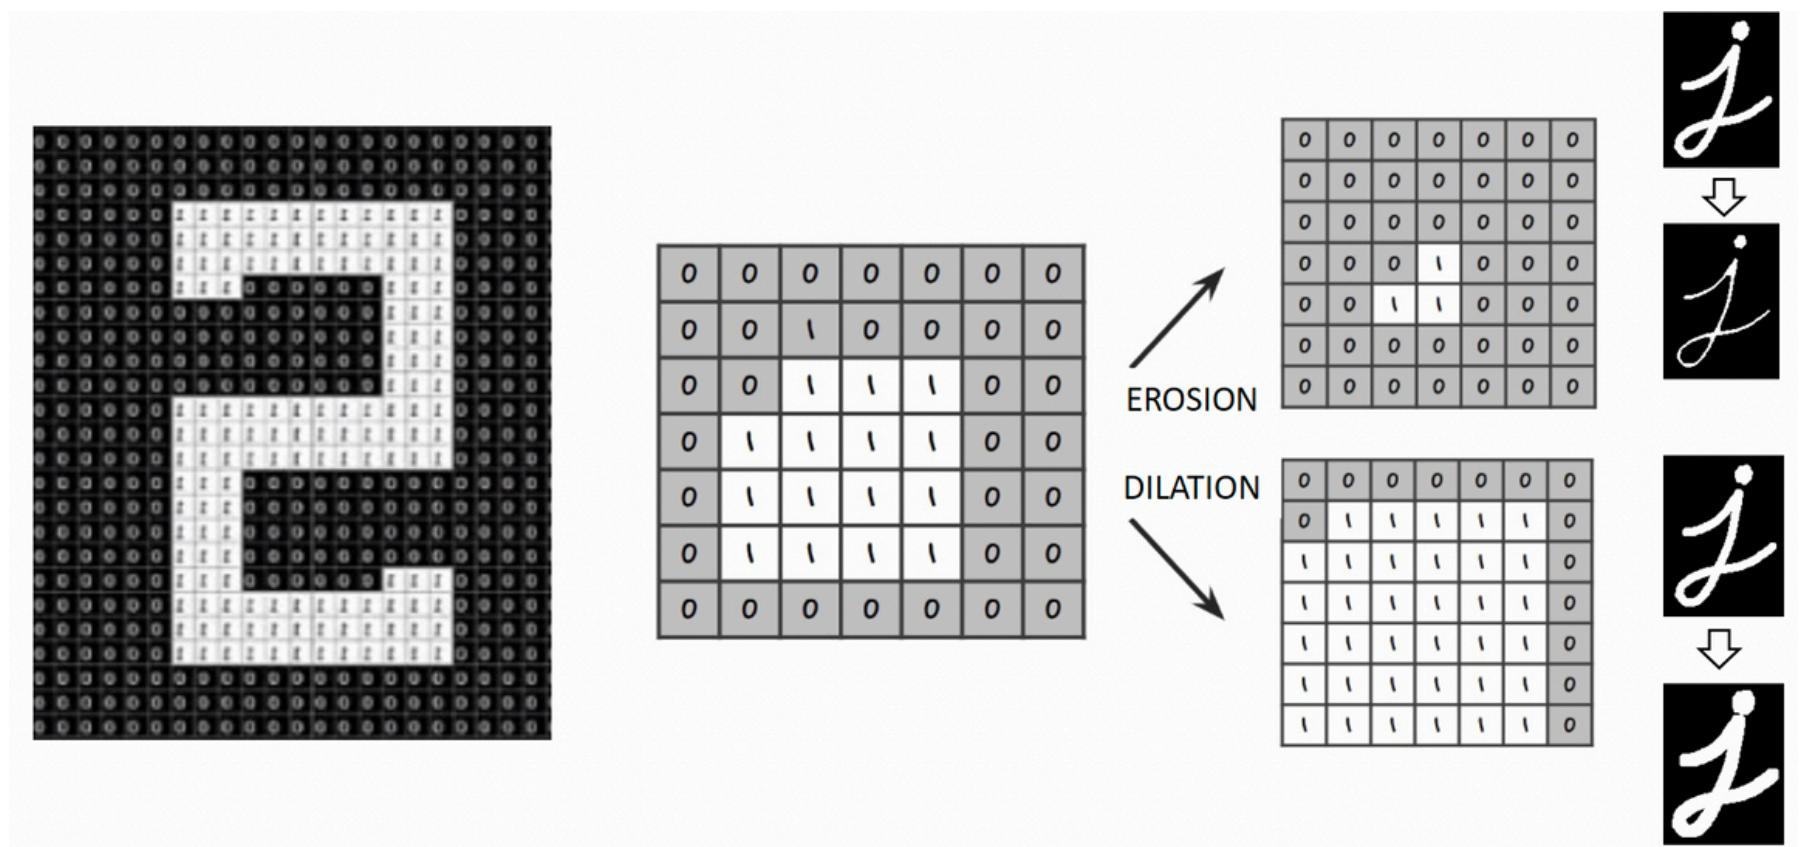


Morphological Operations in OpenCV

1. Erosion

Definition. : Removes pixels at the boundaries of objects in an image.

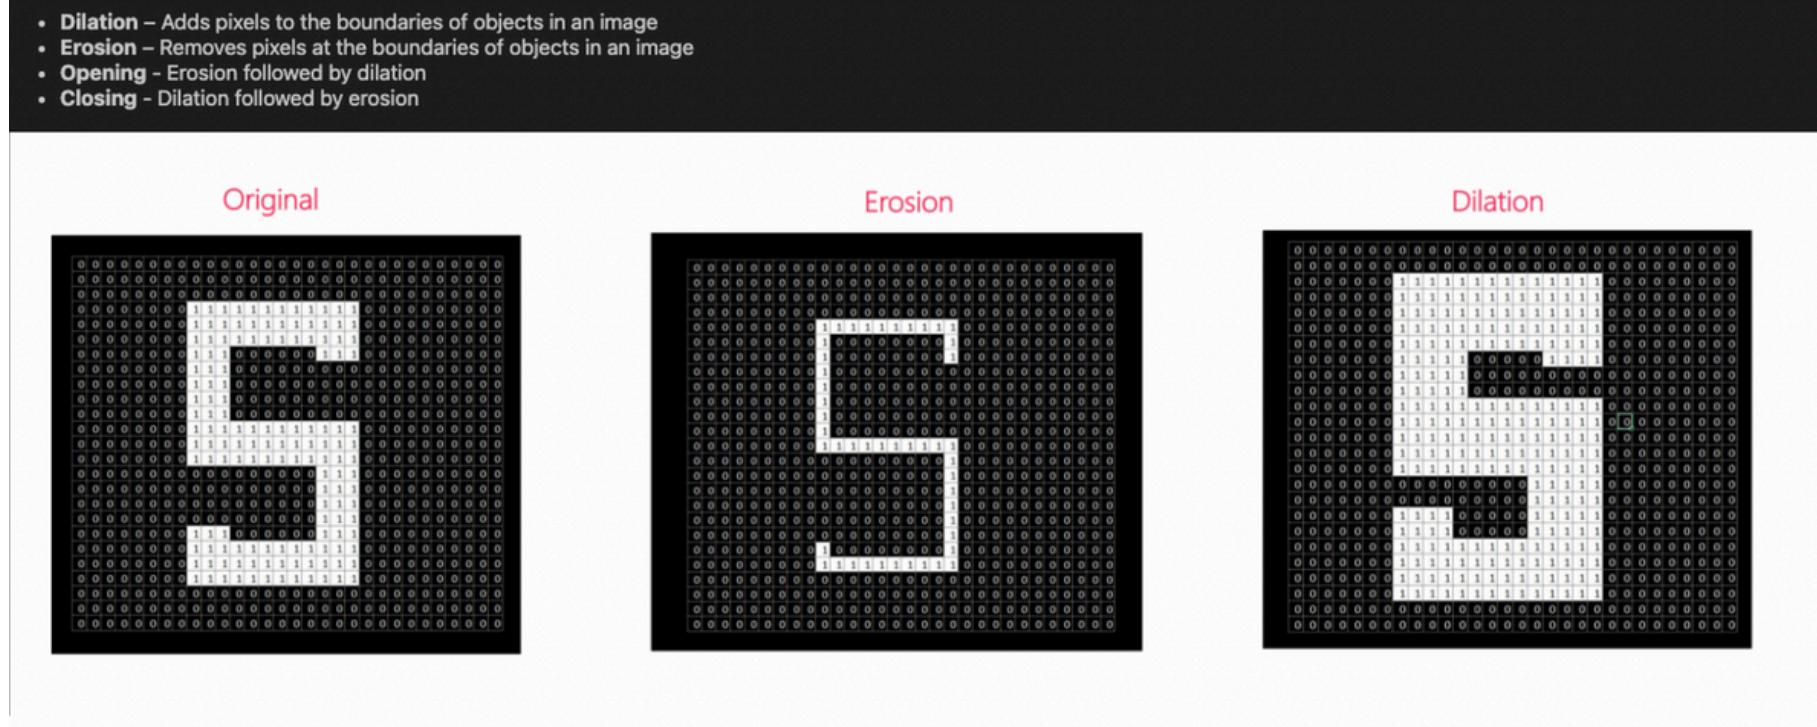
Applications: Used to remove small noise, separate connected objects, and refine boundaries.



2. Dilation

Definition : Adds pixels to the boundaries of objects in an image.

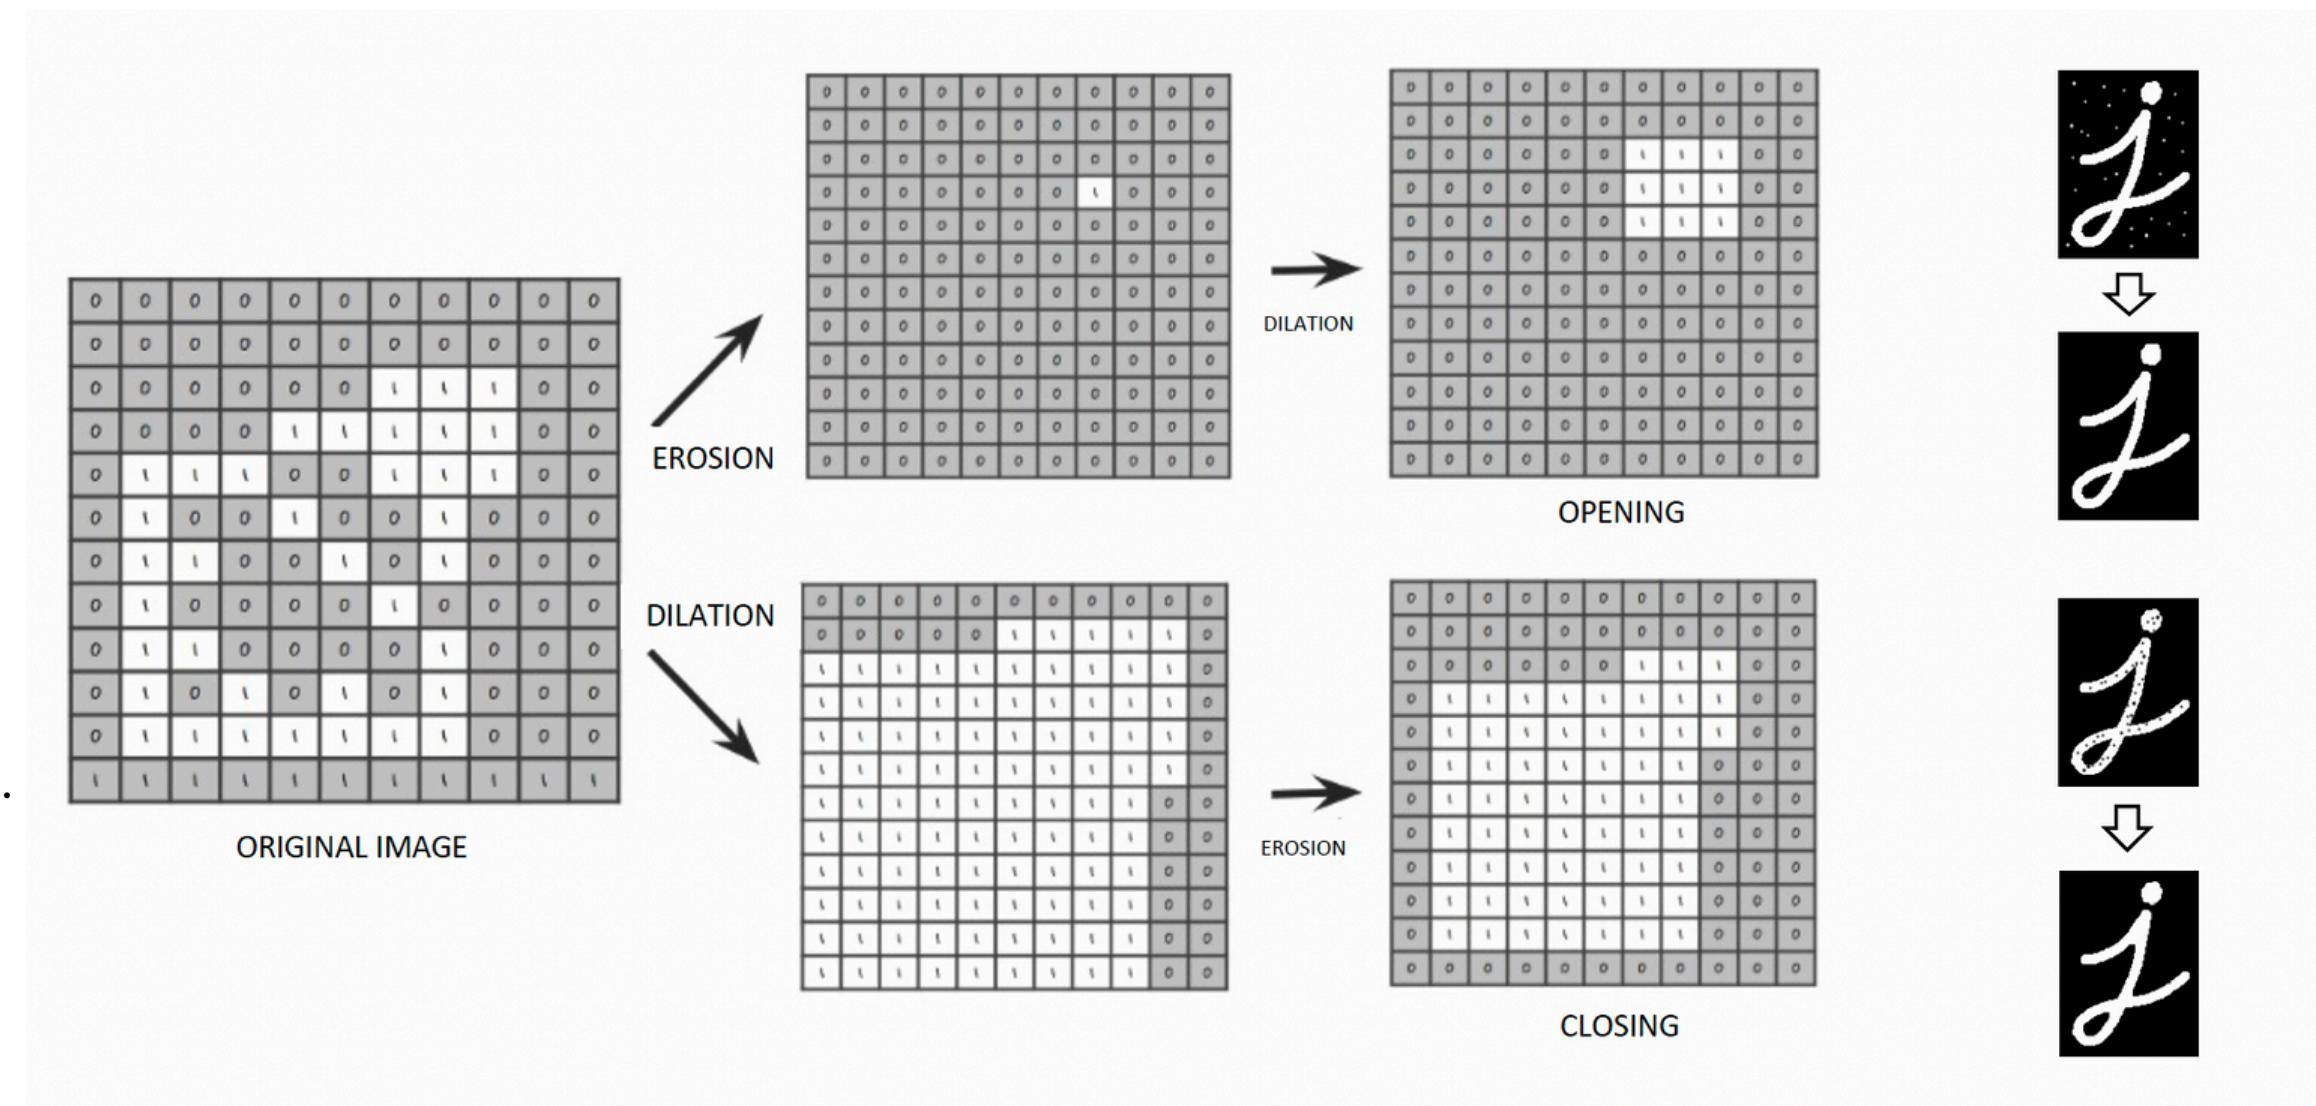
Applications: Useful for increasing object size, reducing noise, and filling small gaps.



Morphological Operations in OpenCV

3. Opening

- Process: Erosion followed by dilation.
- Effect: Removes small objects (noise) from the foreground while keeping the shape and size of larger objects intact.
- Applications: Useful for cleaning up noise in binary images.



4. Closing

- Process: Dilation followed by erosion.
- Effect: Fills small holes and gaps within the objects while maintaining the outer boundary.
- Applications: Often used to close small holes or gaps in the foreground objects.

```
# Define kernels
kernel_3x3 = cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (3 , 3))
# 1. Erosion
eroded = cv2.erode(original, kernel_3x3, iterations=3)
dilated = cv2.dilate(original, kernel_3x3, iterations=3)
```

Original Image



Eroded Image

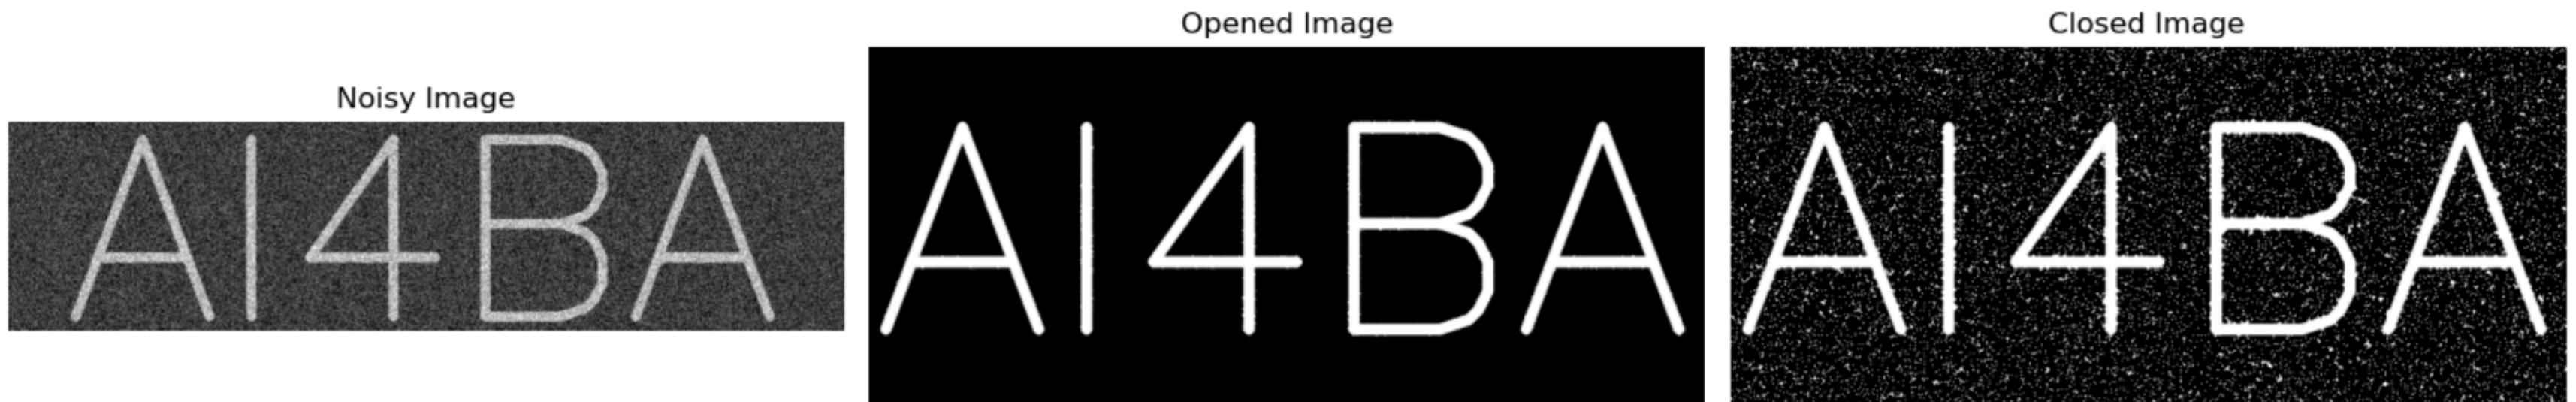


Dilated Image



```
# 3. Opening (Erosion followed by Dilation)
opened = cv2.morphologyEx(result_image, cv2.MORPH_OPEN, kernel_3x3)

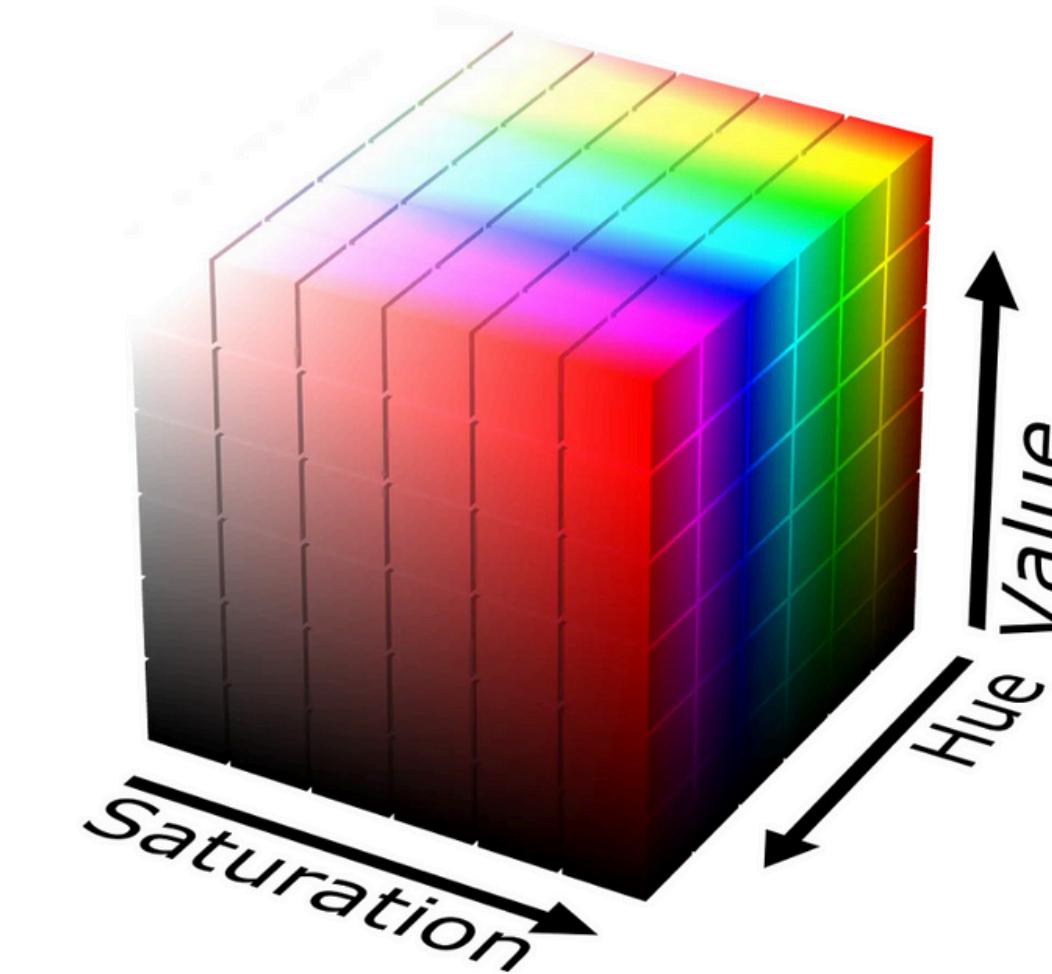
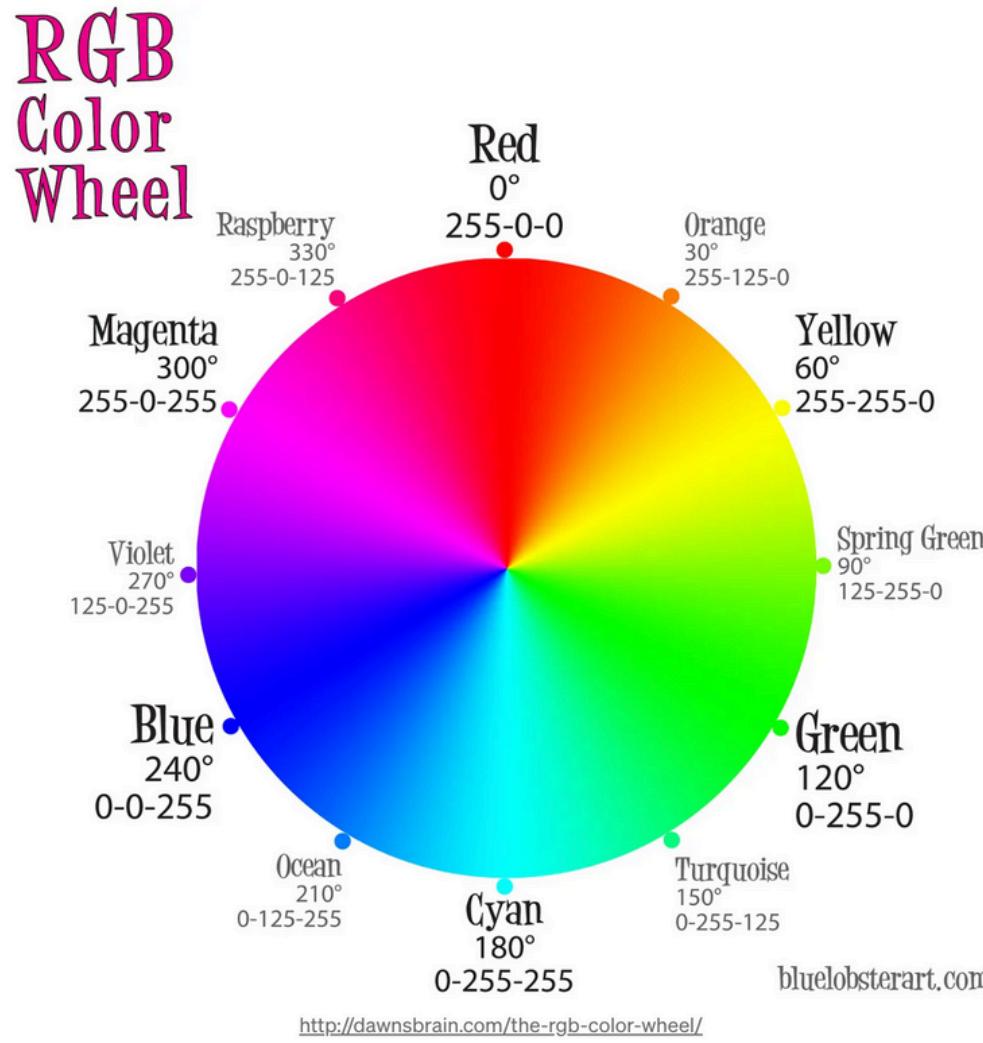
# 4. Closing (Dilation followed by Erosion)
closed = cv2.morphologyEx(result_image, cv2.MORPH_CLOSE, kernel_3x3)
```



Color Spaces and Applications

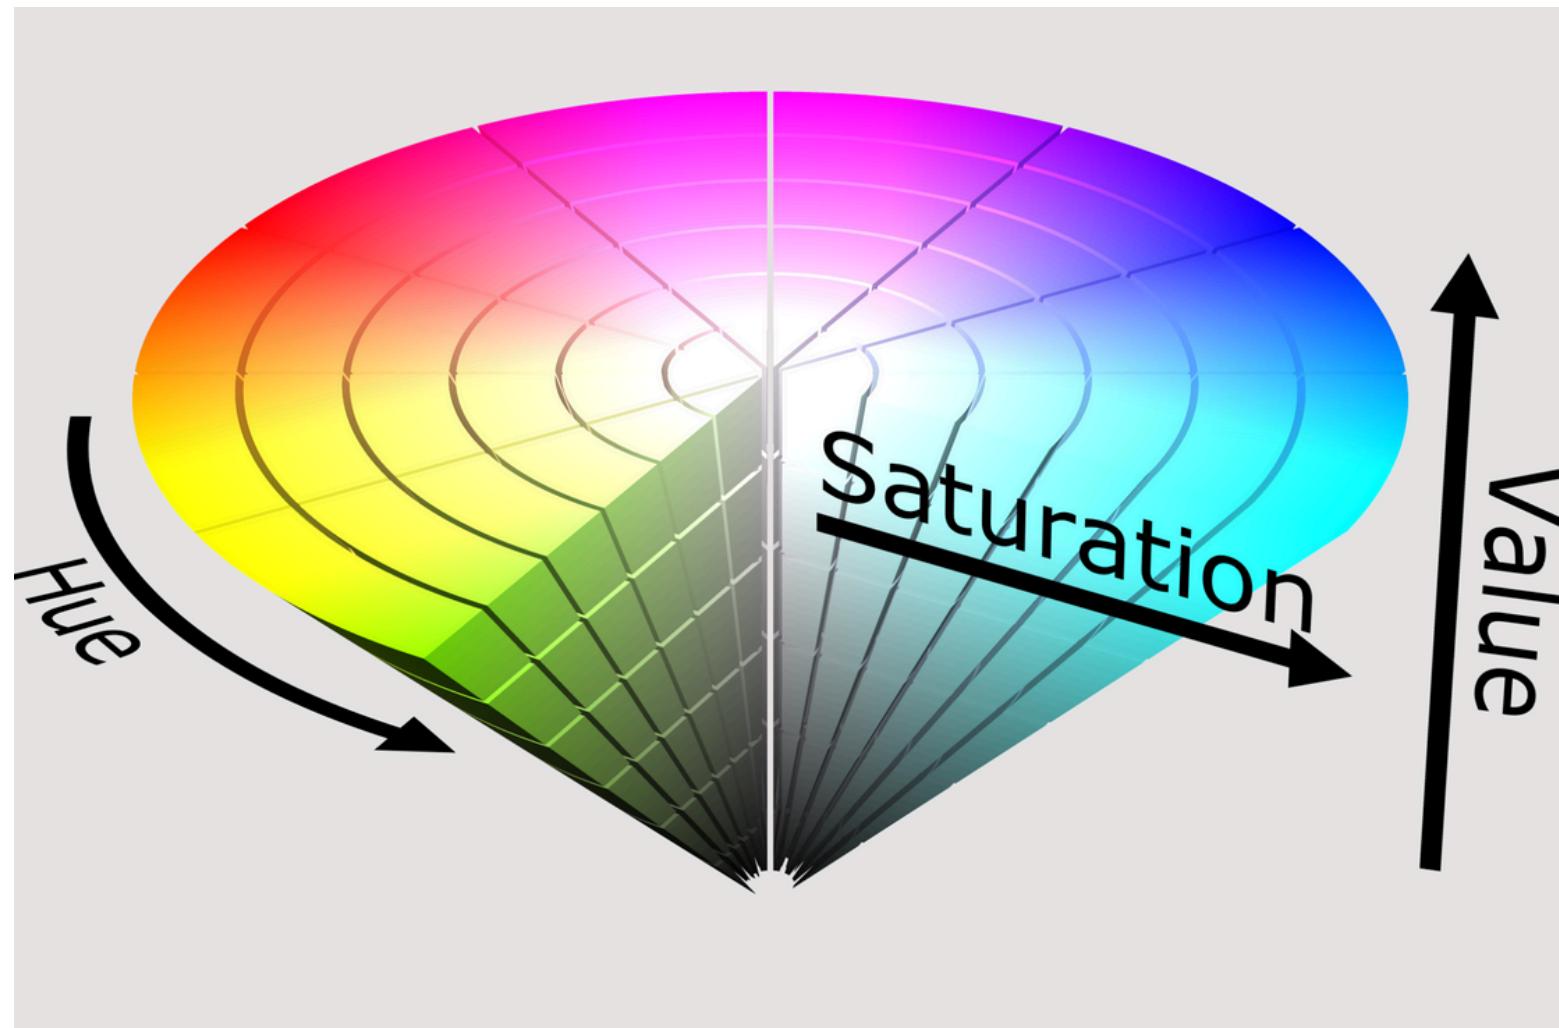
RGB — HSV

RGB เป็นสีที่เกิดจากการผสมกันของสี(แสง) 3 channels (red, green, blue) โดยเราจะเรียกค่าสีที่เปลี่ยนไปตามวงล้อด้านล่างว่า “Hue”



การทำความเข้าใจ Hue ช่วยให้เราสามารถแยกแยะและเลือกใช้สีได้อย่างมีประสิทธิภาพในการออกแบบและงานศิลปะ โดย Hue จะระบุถึงตำแหน่งของสีบนวงล้อสี ซึ่งเป็นเครื่องมือที่ใช้ในการจัดระเบียบและแสดงความสัมพันธ์ระหว่างสีต่าง ๆ

RGB — HSV



ใน OpenCV ปริภูมิสี HSV (Hue, Saturation, Value) มีช่วงค่า

Hue (H): 0 ถึง 179, 0 ถึง 360
Saturation (S): 0 ถึง 255
Value (V): 0 ถึง 255

ช่วงค่าของ Hue ถูกกำหนดให้มีค่าตั้งแต่ 0 ถึง 179 แทนที่จะเป็น 0 ถึง 360 เพื่อประยุกต์ให้เข้ากับความจำและเพิ่มประสิทธิภาพในการประมวลผล ดังนั้น ค่าของ Hue ใน OpenCV จะมีความละเอียดครึ่งหนึ่งของค่าปกติ



HSV (Hue, Saturation, Value) เป็นรูปแบบการแทนสีที่ใช้ในคอมพิวเตอร์กราฟิกและการประมวลผลภาพ เพื่อให้การเลือกและปรับแต่งสีเป็นไปอย่างสะดวกและสอดคล้องกับการรับรู้สีของมนุษย์ ปริภูมินี้ประกอบด้วยสามองค์ประกอบหลัก:

Hue (เวดสี): แทนบุณฑูตแคนแนลตั้งของวงกลมสี โดยเริ่มจากสีแดงที่ 0° ผ่านสีเขียวที่ 120° และสีน้ำเงินที่ 240° จนกลับมาที่สีแดงที่ 360°

Saturation (ความอิ่มตัว): แทนระดับห่างจากแกนกลางของวงกลมสี โดยค่าตั้่มหายถึงสีที่ซัดหรือเป็นสีเทา และค่าสูงหมายถึงสีที่สดใส

Value (ความสว่าง): แทนตำแหน่งตามแกนแนวตั้งของวงกลมสี โดยค่าตั้่มหายถึงสีที่มืดหรือดำ และค่าสูงหมายถึงสีที่สว่างหรือขาว

`cv2.cvtColor` is a function in OpenCV used to convert images from one color space to another. It supports various color conversions such as RGB to grayscale, RGB to HSV, and many more.

```
python คัดลอกโค้ด

cv2.cvtColor(src, code, dstCn=0)
```

Parameters:

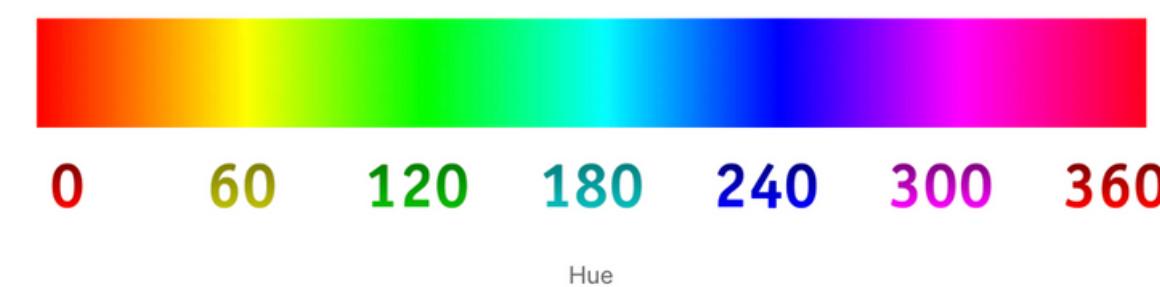
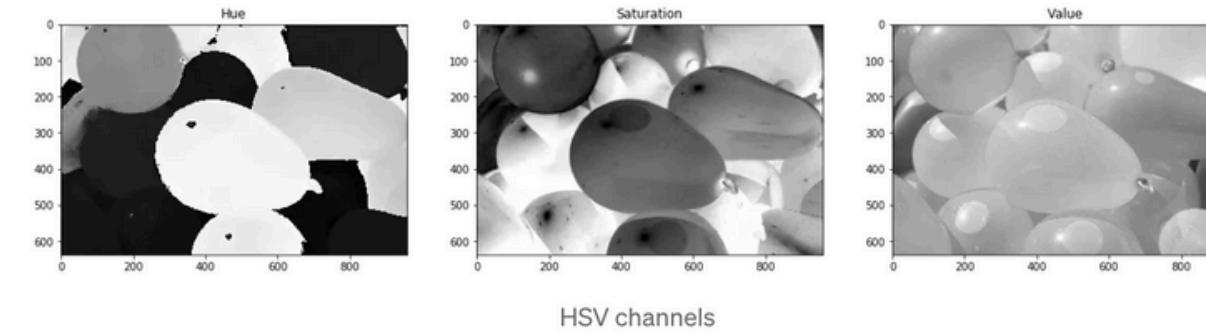
- `src` : The source image to be converted. It should be an image array.
- `code` : The color conversion code. Commonly used codes include:
 - `cv2.COLOR_BGR2GRAY` : Converts BGR to Grayscale.
 - `cv2.COLOR_BGR2RGB` : Converts BGR to RGB.
 - `cv2.COLOR_BGR2HSV` : Converts BGR to HSV.
 - `cv2.COLOR_RGB2BGR` : Converts RGB to BGR.
 - `cv2.COLOR_RGB2GRAY` : Converts RGB to Grayscale.
 - `cv2.COLOR_BGR2Lab` : Converts BGR to Lab color space.
- `dstCn` (optional): The number of channels in the destination image. Default is 0, meaning the number of channels will be derived from the conversion code.

It is very easy to segment a given color if we select the range of Hue that we are interested, for example, the pink:



First we need to convert to HSV color space, this can be done with openCV:

```
cv2.cvtColor(image, cv2.COLOR_RGB2HSV)
```



```

import cv2
import matplotlib.pyplot as plt

# Load an image
image = cv2.imread('./images/lisa.jpg')

# Convert BGR to RGB
image_rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

# Convert BGR to Grayscale
image_gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)

# Convert BGR to HSV
image_hsv = cv2.cvtColor(image, cv2.COLOR_BGR2HSV)

# Display the images
plt.figure(figsize=(12, 8))

plt.subplot(2, 2, 1)
plt.title("Original (BGR)")
plt.imshow(image) # Convert BGR to RGB for correct display in matplotlib
plt.axis('off')

plt.subplot(2, 2, 2)
plt.title("RGB")
plt.imshow(image_rgb)
plt.axis('off')

plt.subplot(2, 2, 3)
plt.title("Grayscale")
plt.imshow(image_gray, cmap='gray')
plt.axis('off')

plt.subplot(2, 2, 4)
plt.title("HSV")
plt.imshow(image_hsv)
plt.axis('off')

plt.tight_layout()
plt.show()

```

Original (BGR)



RGB



Grayscale



HSV



0 60 120 180 240 300 360

Hue

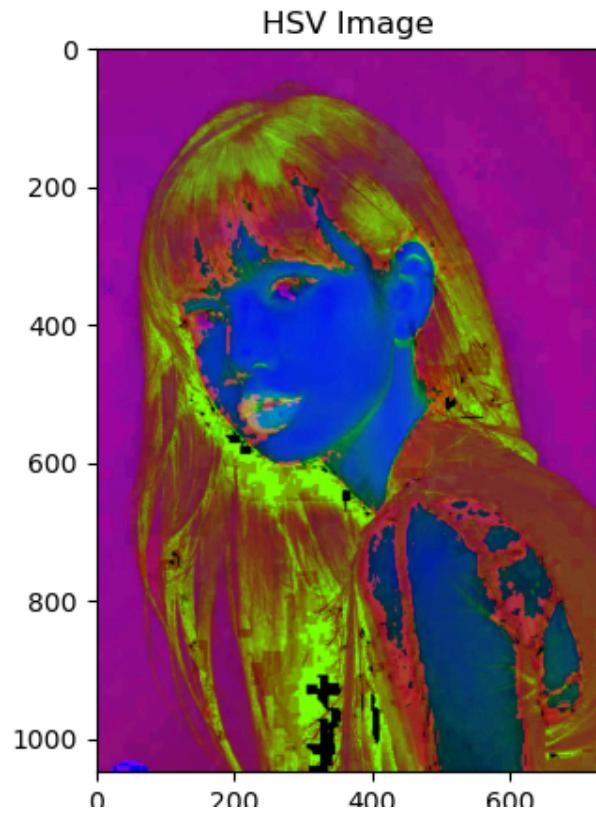
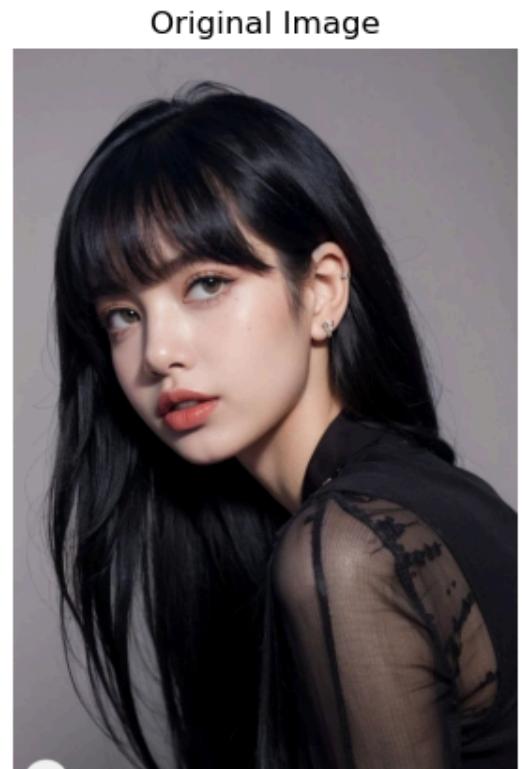
```

import cv2
import numpy as np
from matplotlib import pyplot as plt

# Load the image
image_path = './images/lisa.jpg'
image = cv2.imread(image_path)
image_rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

#Convert to HSV for better color segmentation
hsv = cv2.cvtColor(image, cv2.COLOR_BGR2HSV)
✓ 0.3s

```



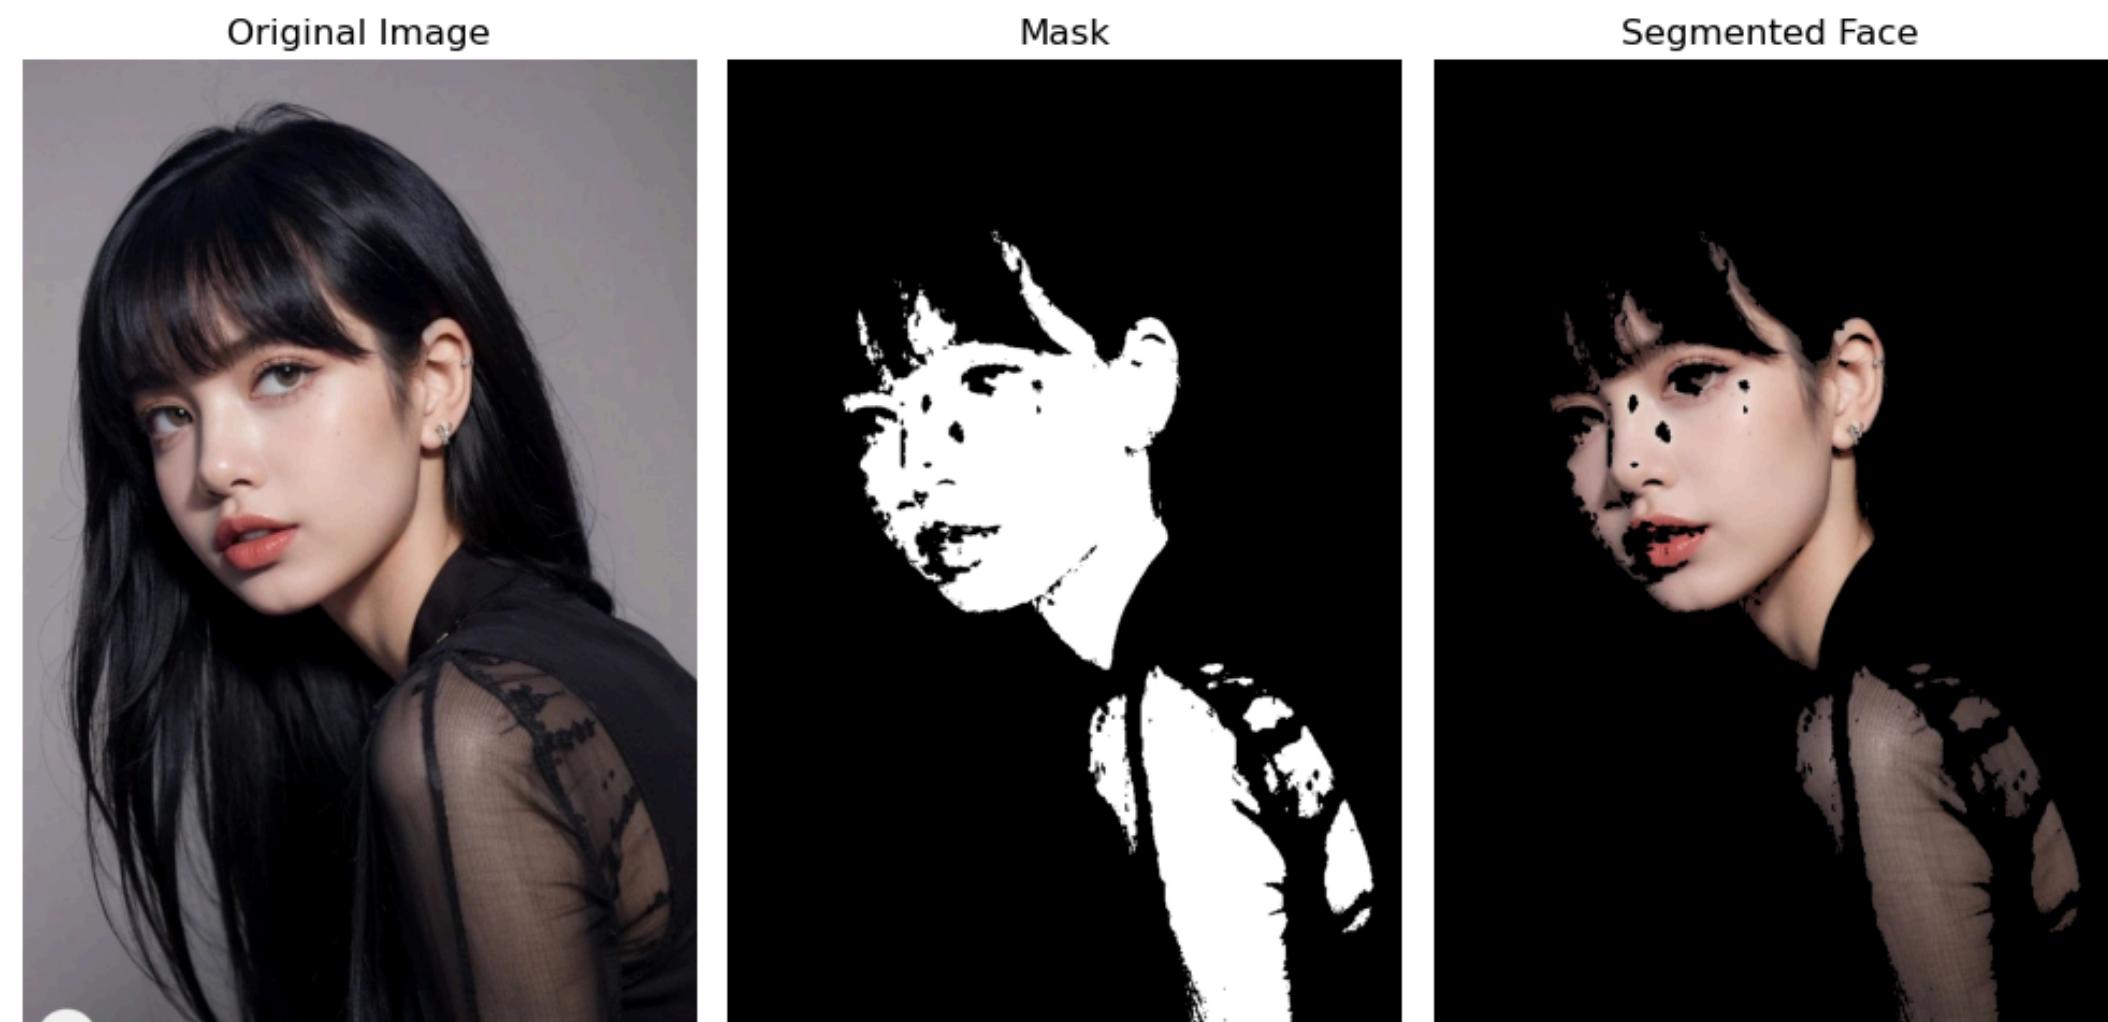
```

# Define the skin color range in HSV
lower_skin = np.array([0, 30, 60], dtype=np.uint8)
upper_skin = np.array([20, 150, 255], dtype=np.uint8)

# Create a binary mask where skin color matches
mask = cv2.inRange(hsv, lower_skin, upper_skin)

#Bitwise AND to extract the face region
result = cv2.bitwise_and(image_rgb, image_rgb, mask=mask)

```



```
# Apply morphological operations to clean the mask
kernel      = cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (5 , 5))
mask_cleaned = cv2.morphologyEx(mask, cv2.MORPH_CLOSE, kernel, iterations=10)

# Bitwise AND to extract the face region
result      = cv2.bitwise_and(image_rgb, image_rgb, mask=mask_cleaned)
```

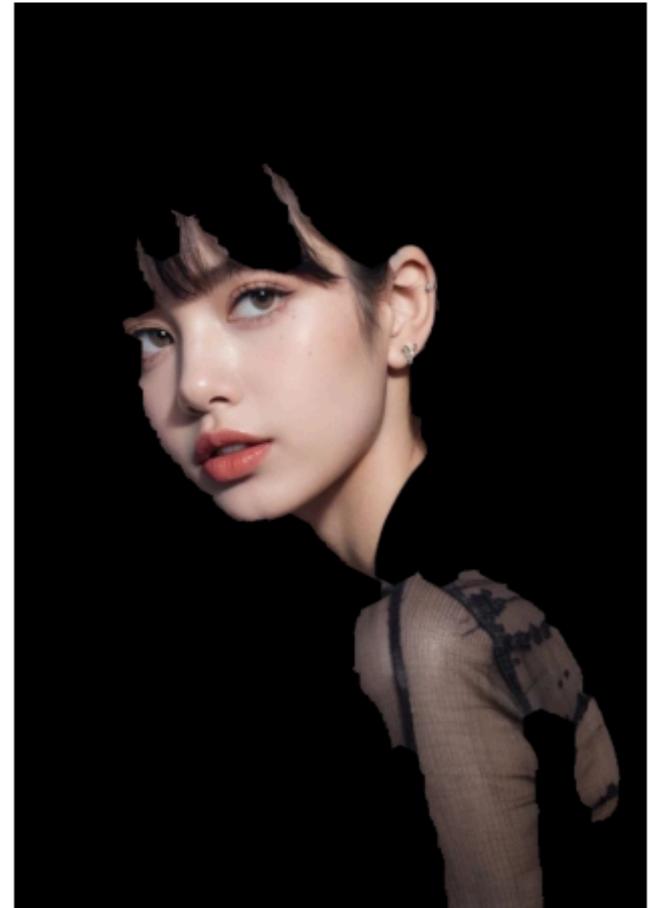
Original Image



Mask



Segmented Face



before morphological



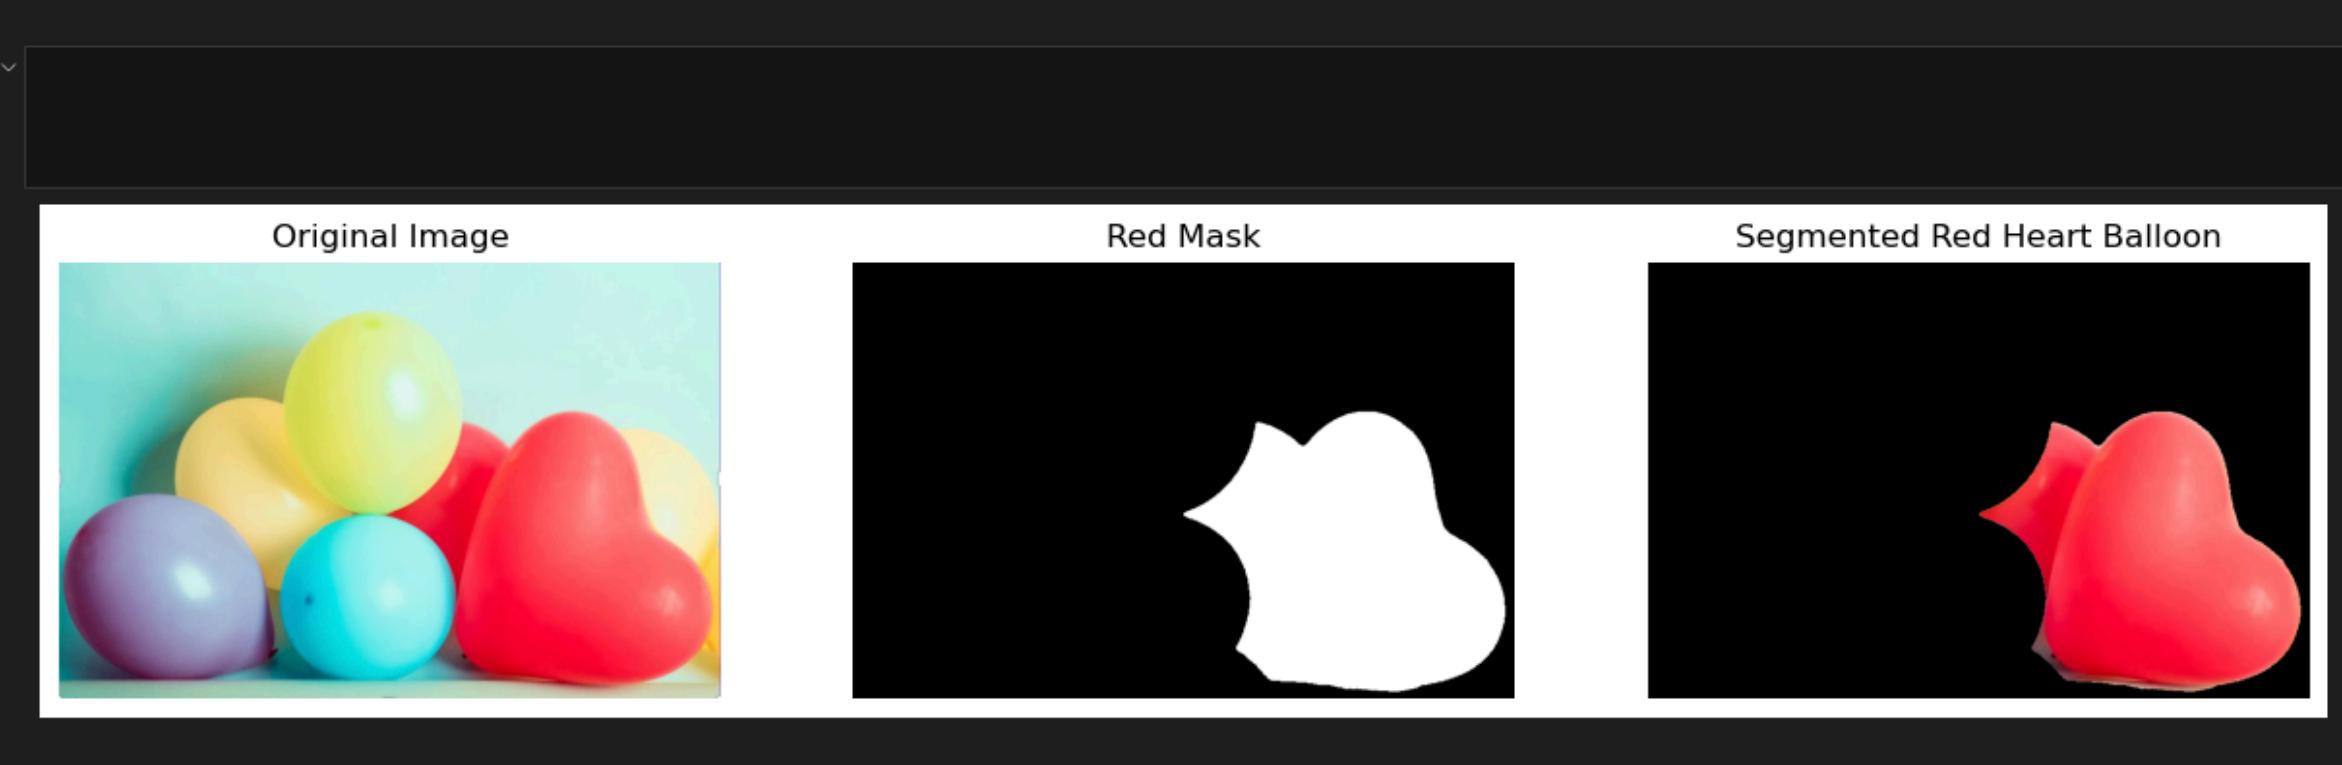
after morphological



Homework -> ให้ดึงข้อความ ฝนทั้งให้เป็นเข็ม ออกแบบตามรูป



Homework ---> ให้ Segment ballon หัวใจ



Homework ---> ໃຫ້ Segment cracked egg

Press `*` `I` to ask GitHub Copilot to do something. Start typing to dismiss.

`<matplotlib.image.AxesImage at 0x136c64220>`

